

## Tutorial 5 Solutions


1. Note that in both 2-4 and B-trees, the heights of all the children of any node are the same. With this requirement, would 2-6 trees or 5-7 trees make sense. What about general A-B trees?

Let us look at an A-B tree, that is, an internal node, except the root can have between A and B children, both inclusive. Let us see what happens during:

- Delete: A node x may then have A-1 children and become inconsistent. Then the first option is to check if one of its siblings' children may be transferred. That can be done if either has more than A children. If both have A, then there is no option but to merge x with a sibling y to a new node z. Then z will have  $A + A - 1 = 2A - 1$  children. Thus we must have  $2A - 1 \leq B$ , for z to be consistent.
- Insert: Suppose a node x becomes inconsistent due to its children going from B to B+1. It could shed one of its children with its siblings. But what if both have B children? Then we must split x into 2 nodes, each with at least A children. Thus we must have  $B - 1 \geq 2A$ , giving us the same condition. Thus 2-4 trees make sense but not 5-7.

2. Show two examples of 2-4 trees where an insertion increases the levels by 1 and a deletion drops the level by 1. Can these be the same trees?

See the examples below. No, the two trees cannot be the same since they must differ at the root level. While a tree which causes the level to increase must have B children at its root, a tree which causes the level to drop will have 2 children. Note that this does not depend on A or B.



3. Why is insertion in 2-4 and B-trees simpler than deletion?

This is because an insert must happen at the leaf level. A delete can be at any internal node. This makes all the difference. If deletions were at the leaf nodes then delete would not be as complicated.

4. In real-life B-trees, the number of children could be 100-1000. In this case, what is the fraction of values which are stored at the leaf level of the total number of values?

For an A-B tree of L levels, the minimum number of keys are  $(A-1) + (A-1)*A + (A-1)*A^{L-1} + A^L$ . The maximum number is the same replaced by B. Thus, the number of internal keys are  $(B-1)(B^L-1)/(B-1) = B^L - 1$  which is of the same order as the last level! So really, even for large Bs, storing just at the leaves would be quite wasteful.

5. Design a structure where the values are stored only at the leaf level. Only end-markers are stored in the intermediate levels. Would insertion and deletion be easier in such "simple-B-trees"?

This is not difficult. For any internal node, with say B children, we store the numbers  $[(L_1, U_1)(L_2, U_2) \dots (L_B, U_B)]$  where  $(L_i, U_i)$  is the lowest and the highest key values stored under that child. It is then easy to delete since it happens at the leaf level.

