

TCP Congestion Control -- Overview

Kameswari Chebrolu

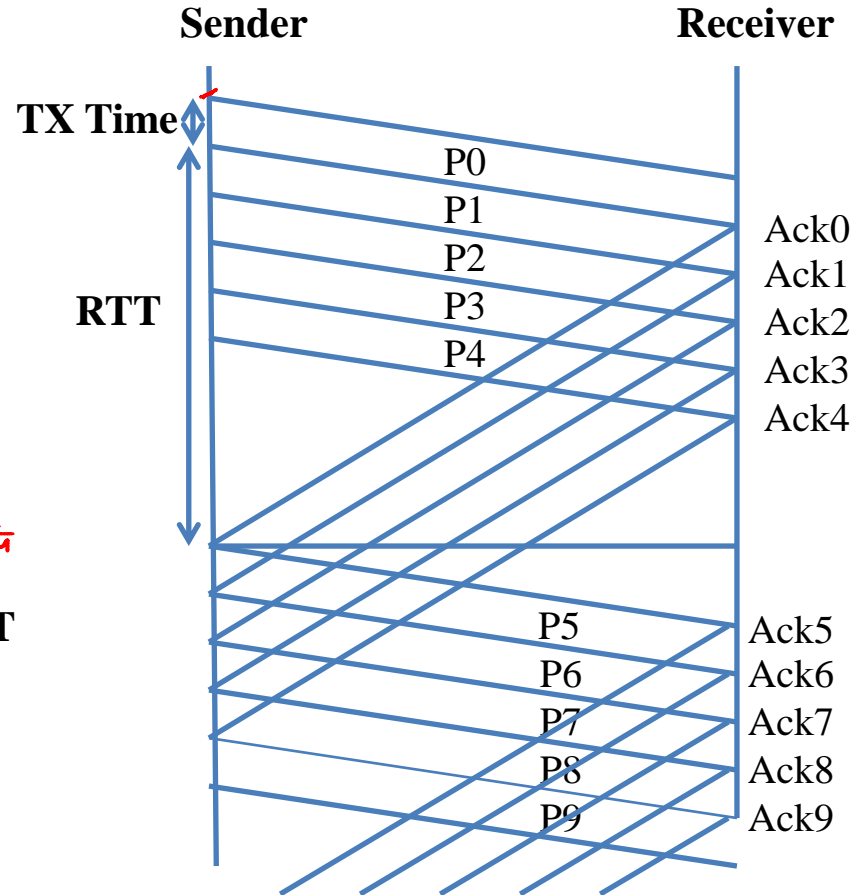
Seminal Paper: Congestion Avoidance and Control
by Van Jacobson and Michael J. Karels

Recap: TCP Services

- Multiplexing/Demultiplexing
- Reliable point-to-point data transfer
- Full-duplex
- Congestion control
- Flow control

Sliding
window
protocol

Recap: Sliding Window




Throughput $\sim (W * MSS) / RTT$

Handwritten annotations:

- A red circle around **W** with an arrow pointing to it and the text "wind. size" (window size).
- A red arrow pointing up to **MSS** with the text "pkts" (packets).

Congestion Control: Problem Statement



- Many flows pass through a router; number varies with time
- Flows can be TCP or UDP ^{fairness}
- The link capacities of the routers are different ^{blocks}
- End Result: Throughput achieved by a given flow function of many factors

- Need to estimate W (of sliding window) such that each flow gets its fair share
 - Estimate small \rightarrow underutilization; Estimate large \rightarrow Congestion
- W will vary over time
- Congestion Control: Preventing sources from sending too much data too fast and thereby ‘congest’ the network

Sliding Window Protocol

- Roughly, idea translates to the following:
- View network as a pipe
- Determine the capacity of the pipe (Bandwidth-delay product)
- Fill the pipe with data
- As you remove one packet from the pipe, add another
 - ACKs help clock out data (Self Clocking)

3 Steps

- Getting to Equilibrium[↗]
- Conservation at equilibrium[→]
 - Don't put new packet unless old one is removed
- Adapting to Path Dynamics

Summary

- Congestion Control is a complex problem
- Need to implement it in the context of the sliding window protocol
 - Self clocking[✓] is a useful feature
 - Need to determine and adapt W (window size) such that you don't underutilize bandwidth or congest the network
- Ahead: Actual details —