

Computer Architecture

Instruction Set Architecture

Virendra Singh

Professor

Computer Architecture and Dependable Systems Lab

Department of Electrical Engineering
Indian Institute of Technology Bombay

<http://www.ee.iitb.ac.in/~viren/>

E-mail: viren@ee.iitb.ac.in

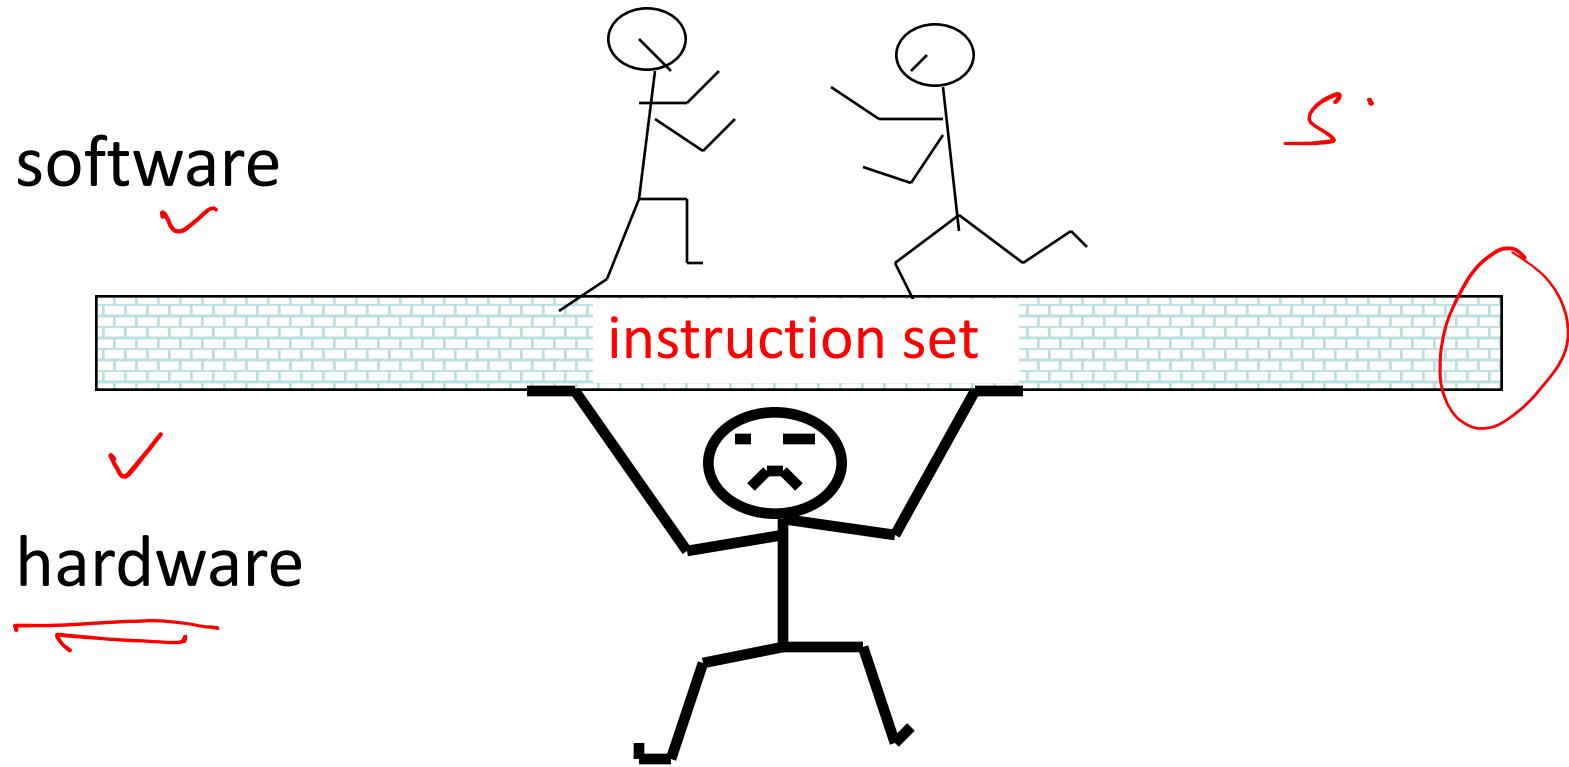
CS-683: Advanced Computer Architecture

Lecture 5 (09 August 2021)

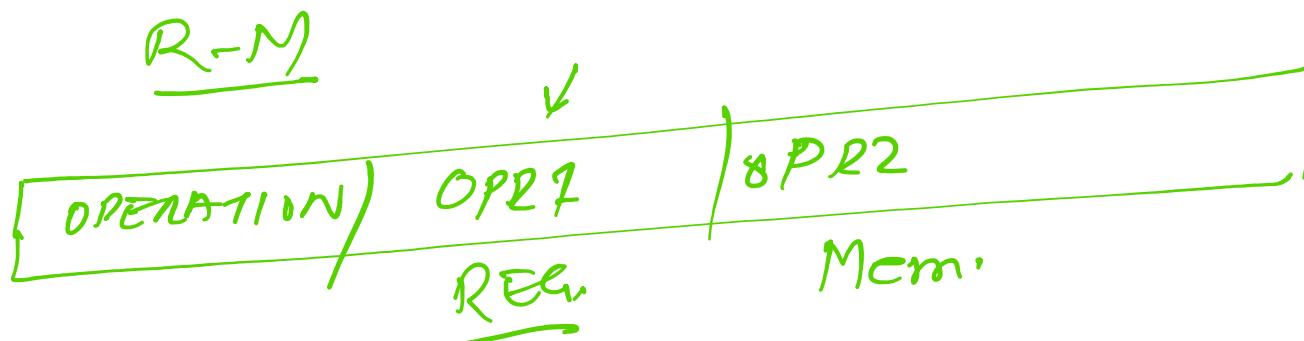
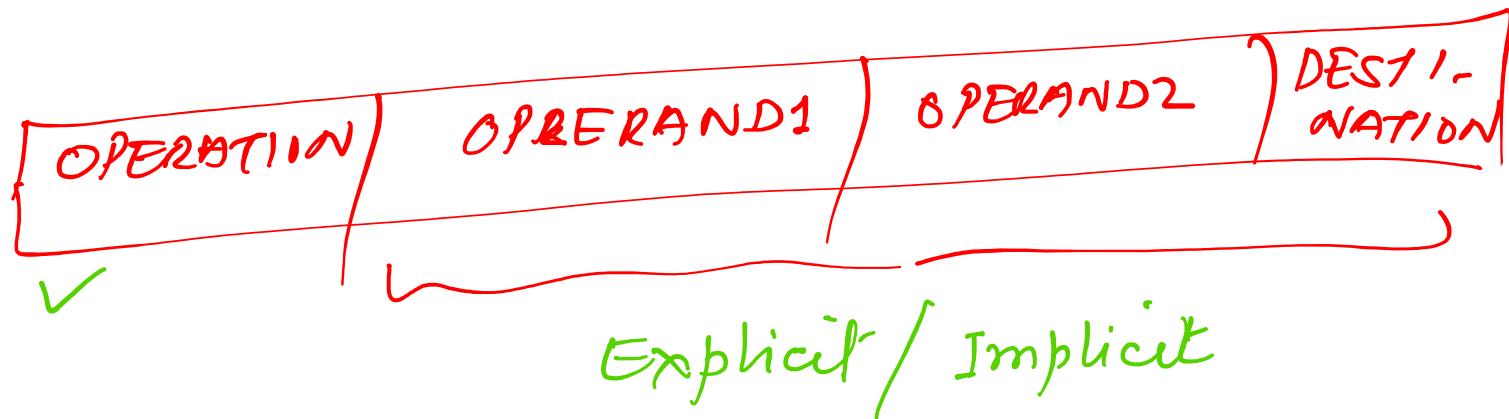


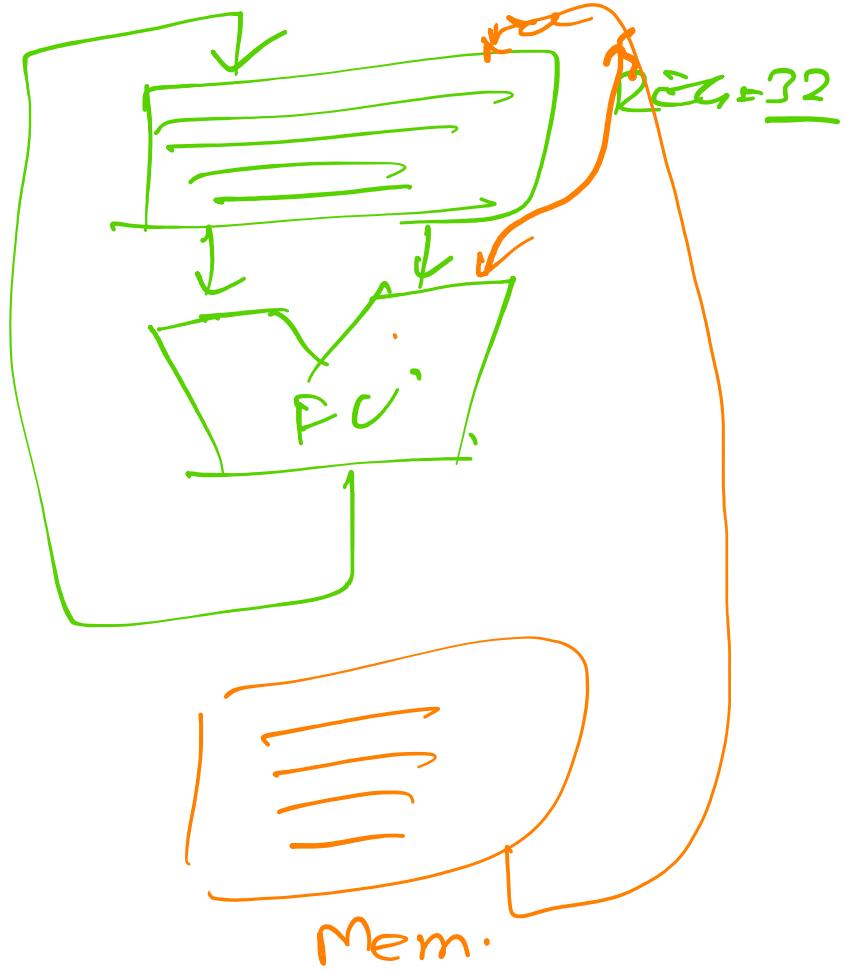
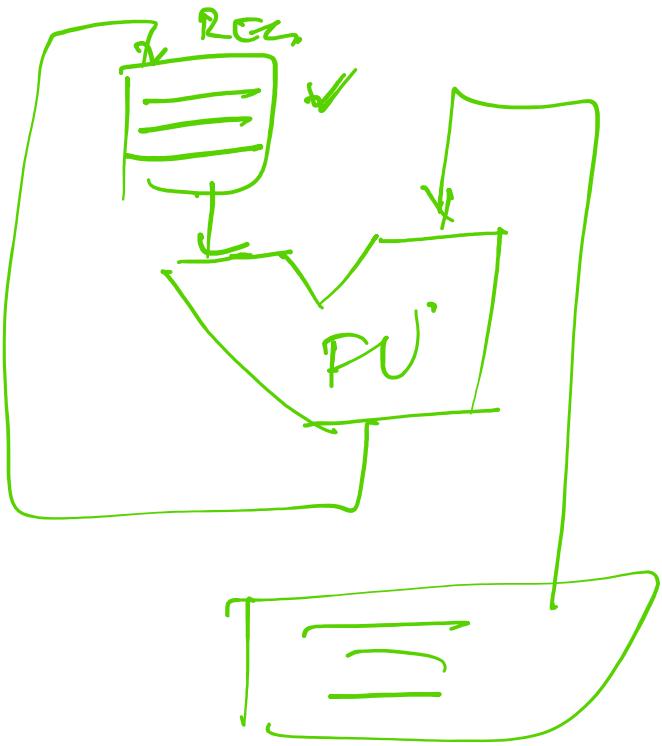
CADSL

Instruction Set Architecture (ISA) ✓



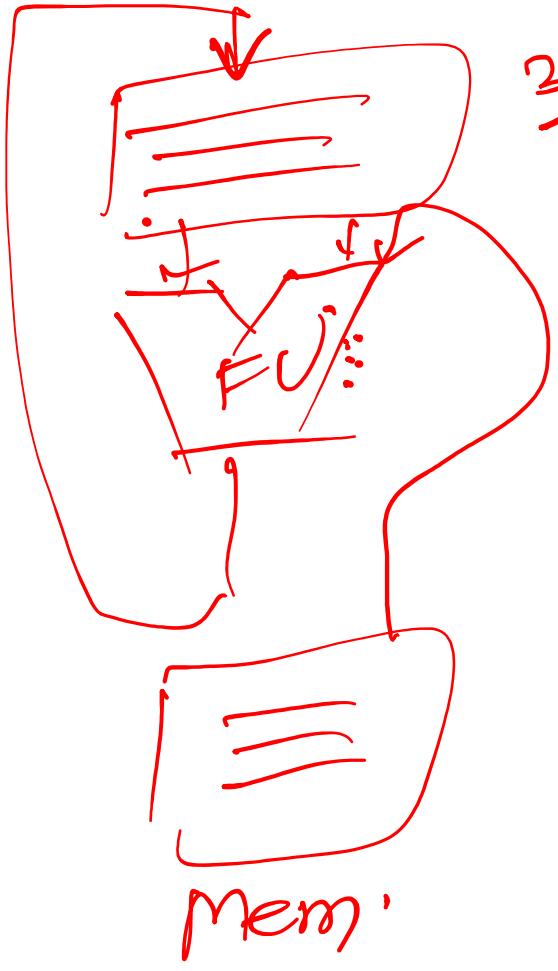
1. Scarcity of memory \Rightarrow min # bits to encode an inst
2. Slow memory
3. Ease of programming



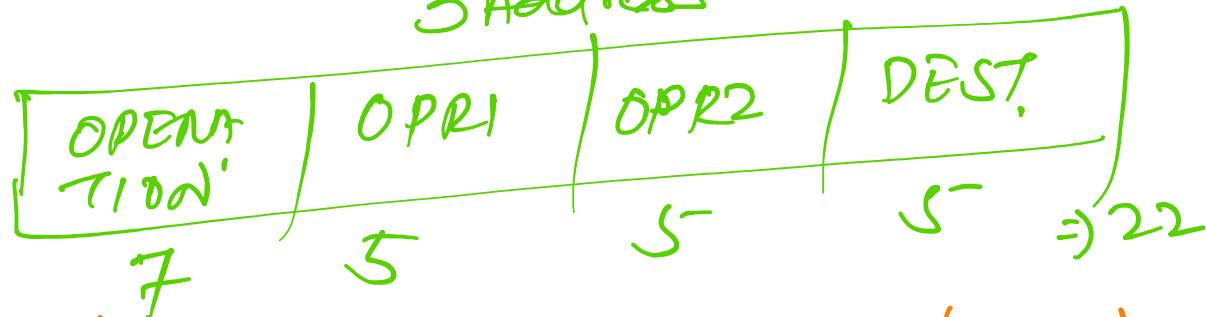
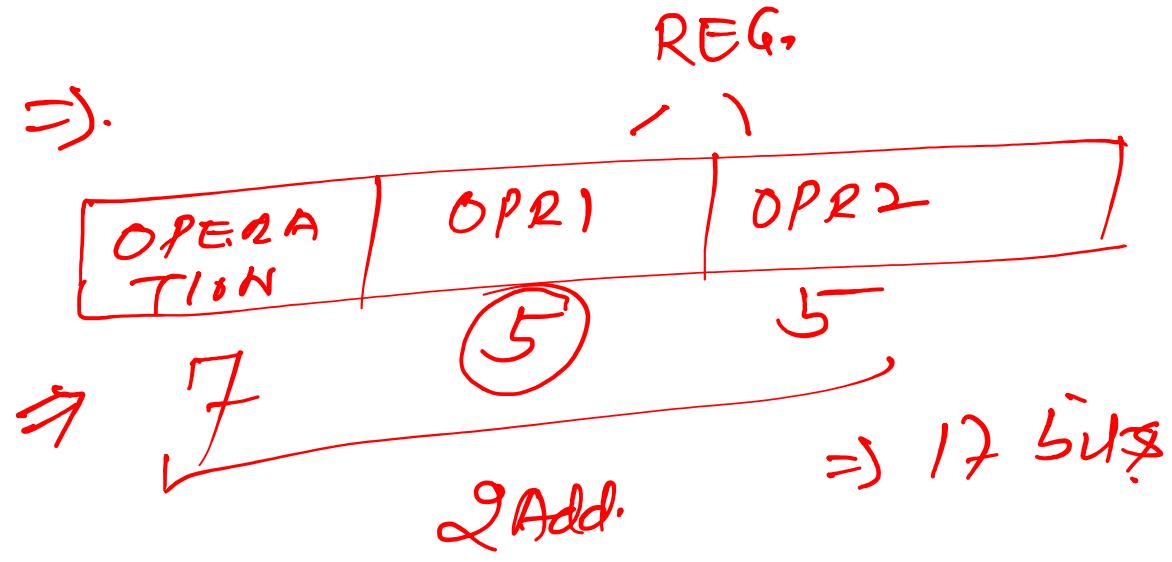


Memory traffic
Slow





32
⇒



Arithmetic / Logic / control flow (branch/jumps)

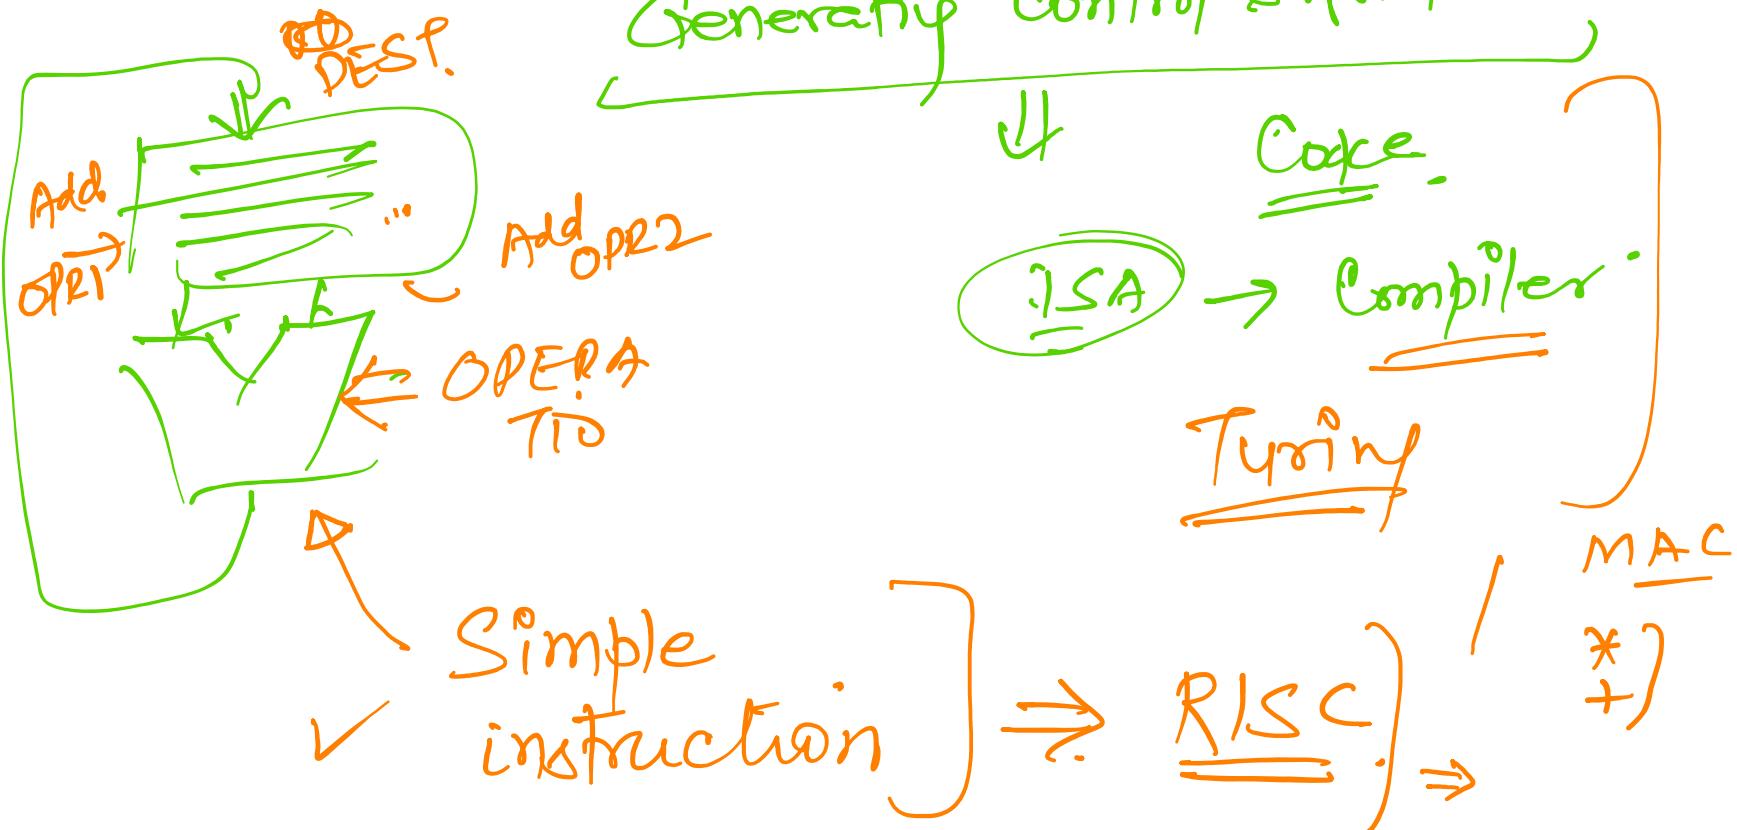
Load / Store ⇒ allowed to communicate with memory,



Load / Store Arch.

very close to machine

Generating control signal for machine.



Load-Store Architectures

- Instruction set:

add R1, R2, R3

sub R1, R2, R3

mul R1, R2, R3

load R1, R4

store R1, R4

- Example: $A^*B - (A+C^*B)$

load R4, A

load R5, B

load R6, C

mul R7, R6, R5

/* C*B */

add R8, R7, R4

/* A + C*B */

mul R9, R4, R5

/* A*B */

sub R10, R9, R8

/* A*B - (A+C*B) */

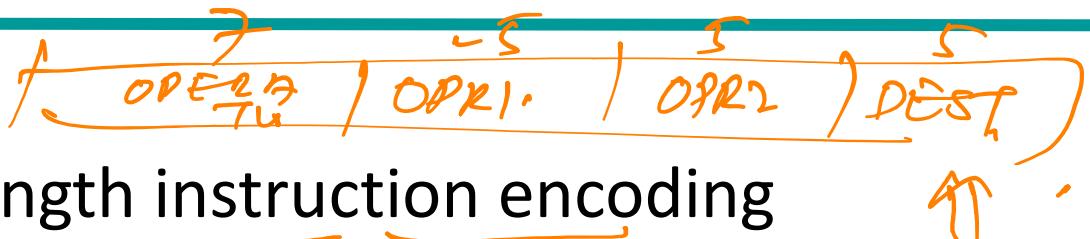
$$\underline{R_1} = \underline{R_2} + \underline{R_3}$$



Load-Store: Pros and Cons

- Pros

- Simple, fixed length instruction encoding
- Instructions take similar number of cycles
- Relatively easy to pipeline



- Cons

- Higher instruction count
- Not all instructions need three operands
- Dependent on good compiler



Registers: Advantages and Disadvantages

- Advantages

- Faster than cache (no addressing mode or tags)
- Deterministic (no misses)
- Can replicate (multiple read ports)
- Short identifier (typically 3 to 8 bits)
- Reduce memory traffic

M - M

- Disadvantages

- Need to save and restore on procedure calls and context switch
- Can't take the address of a register (for pointers)
- Fixed size (can't store strings or structures efficiently)
- Compiler must manage



How Many Registers in RF

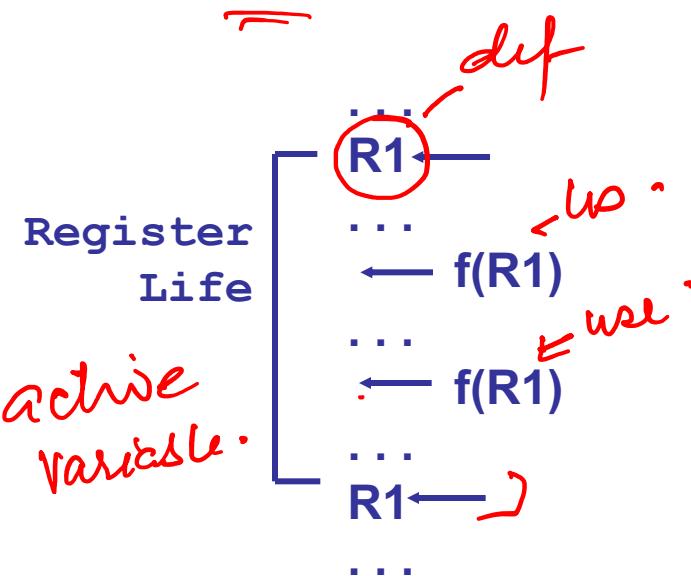
6 algorithms from CALGO (ACM) written in 4 languages;
ALGOL, BASIC, BLISS, FORTRAN

Register Life

Life Length	No. of Lives	Fraction	Cum. Fraction
1~1	174,927	0.09	0.09
2~3	728,346	0.38	0.48
4~7	547,072	0.29	0.77
8~15	252,508	0.13	0.90
16~31	116,404	0.06	0.96
32~63	41,673	0.02	0.98
64~127	17,790	0.01	0.99
128~	15,603	0.01	1.00

Avg number of simultaneous RL: 2 ~ 6

No program uses more than 15 registers simultaneously



We need to try to keep the live registers in the RF

yy.



Slow Memory

CISC \Rightarrow RISC

Memory $200 - 1000$ time slower.

Graph showing time vs. memory access:

\Rightarrow Place as much work as we can.

INSERT A, \rightarrow .

FTI.



Complex vs. Simple Instructions

- **Complex instruction:** An instruction does a lot of work, e.g. many operations
 - Insert in a doubly linked list
 - Compute FFT
 - String copy
- **Simple instruction:** An instruction does small amount of work, it is a primitive using which complex operations can be built
 - Add ✓
 - XOR ✓
 - Multiply ✓



Complex vs. Simple Instructions

- Advantages of Complex instructions

- + Denser encoding → smaller code size → better memory utilization, saves off-chip bandwidth, better cache hit rate (better packing of instructions)
- + Simpler compiler: no need to optimize small instructions as much

$$\begin{aligned}f &= s + b \times c \\&\quad \overbrace{\qquad\qquad\qquad}^{\text{MAC}} \\t_1 &= b \times c \\f &= s + t_1\end{aligned}$$

- Disadvantages of Complex Instructions

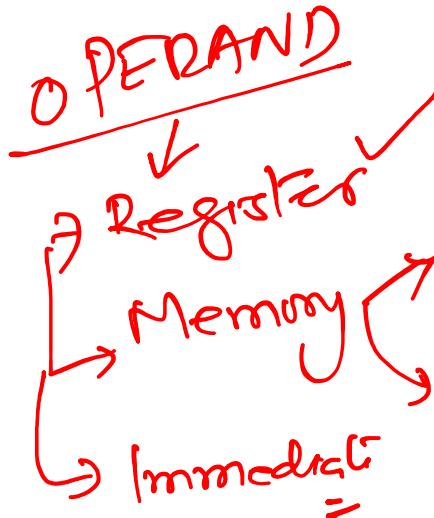
- Larger chunks of work → compiler has less opportunity to optimize (limited in fine-grained optimizations it can do)
- More complex hardware → translation from a high level to control signals and optimization needs to be done by hardware



3

Ease of Programming

Reg.
↓
scalar



for ($i = 0; i < 1000; i++$)
 $A(i) = A(i) + B(i)$

which Address: \downarrow

$A(0) = B(0) + B(0)$
 $A(1) = A(1) + B(1)$
;
 $A(i) = B(i) + B(i)$

Add
 $a_i = b_i + s_i$

TOPDATA (A_{102}) OPRI | OPR.

T100V ✓ ✓



Kinds of Addressing Modes

~~* *~~

OP	Ri	Rj	v
----	----	----	---



memory

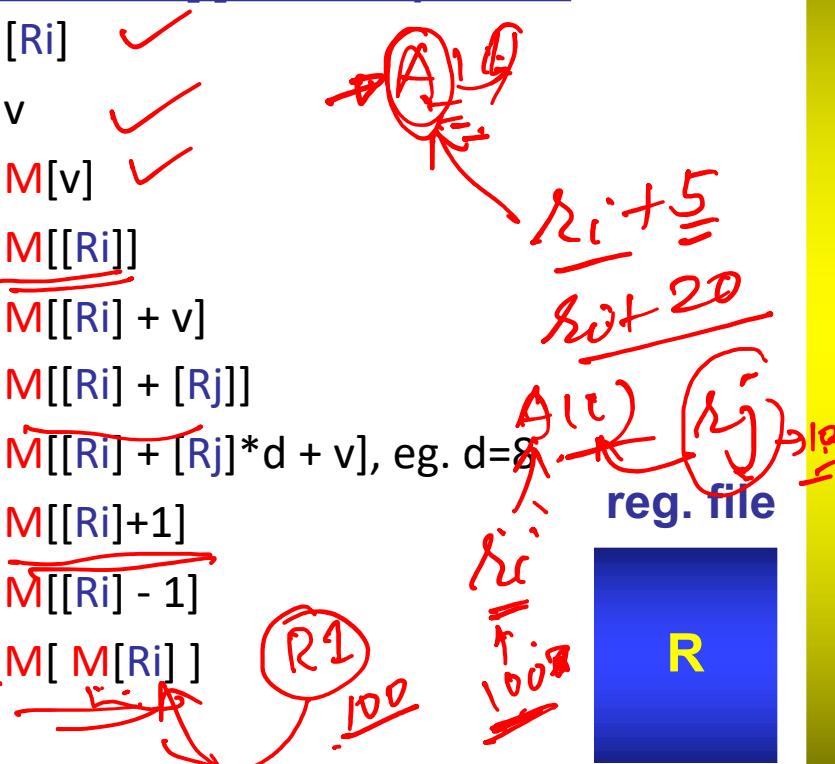


Addressing Mode

- Register direct
- Immediate (literal)
- Direct (absolute)
- Register indirect
- Base+Displacement
- Base+Index
- Scaled Index
- Autoincrement
- Autodecrement
- Memory Indirect
- [Indirection Chains]

value in [] is the operand

- [Ri] ✓
- v ✓
- M[v] ✓
- M[[Ri]]
- M[[Ri]] + v
- M[[Ri]] + [Rj]]
- M[[Ri]] + [Rj]] * d + v, eg. d=8
- M[[Ri]] + 1]
- M[[Ri]] - 1]
- M[M[Ri]]]



100
100
100
50
100
100

09 Aug 2021

CS-683@IITB

14

CADSL

Instruction Execution Characteristics: Type of Operands

Dynamic Frequencies of Occurrences

	PASCAL	C	Average
Integer Constant	16	23	20
Scalar Variable	58	53	55
Array/Structure	26	24	25

✓
a, b, -
A(i)

Majority of references to scalar

- 80% are local to a procedure
- References to arrays/structure require index or pointer

Locations of operands(Average per instruction)

- 0.5 operands in memory
- 1.4 operands in registers]

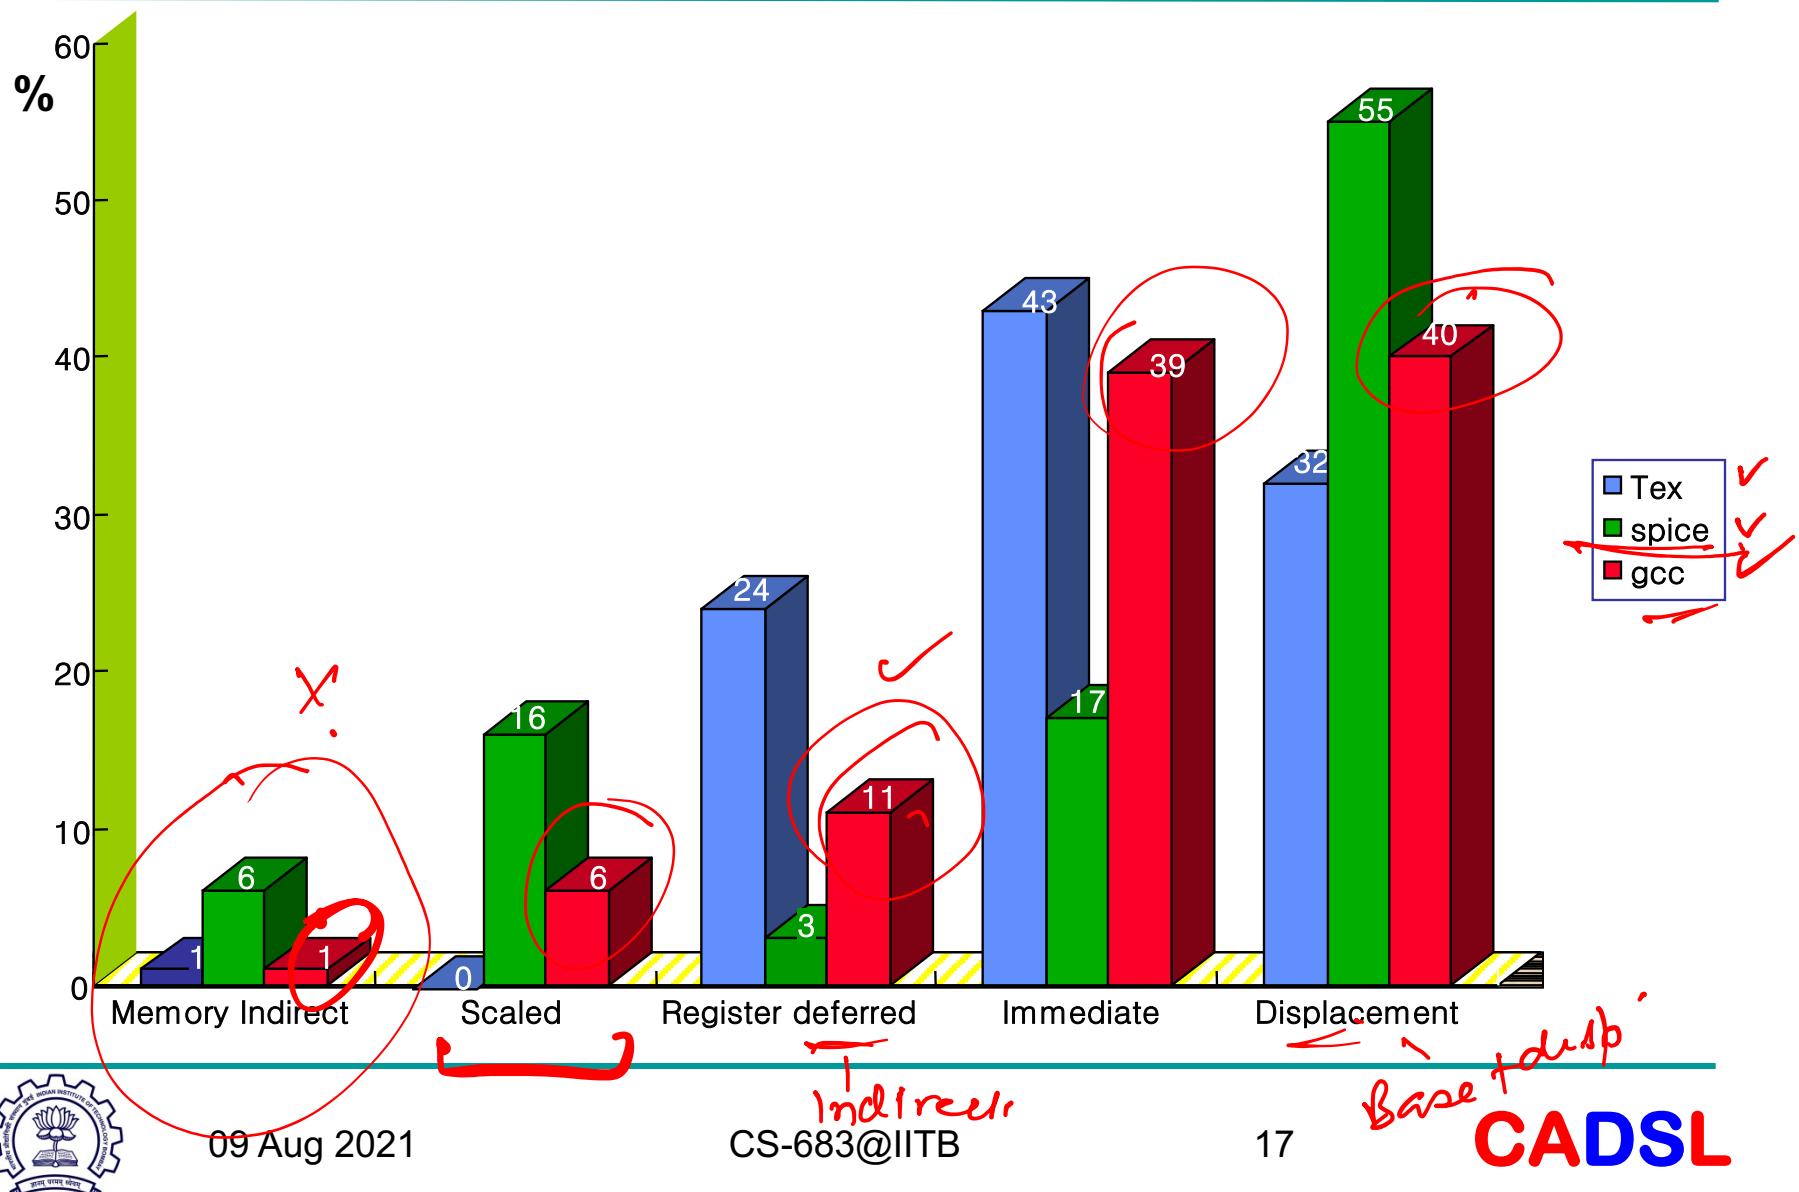


Addressing Modes (VAX)

- Register
- Immediate
- Register Indirect
- Displacement
- Indexed
- Direct Absolute
- Memory Indirect
- Auto Increment
- Auto decrement



Memory Addressing Modes (VAX)



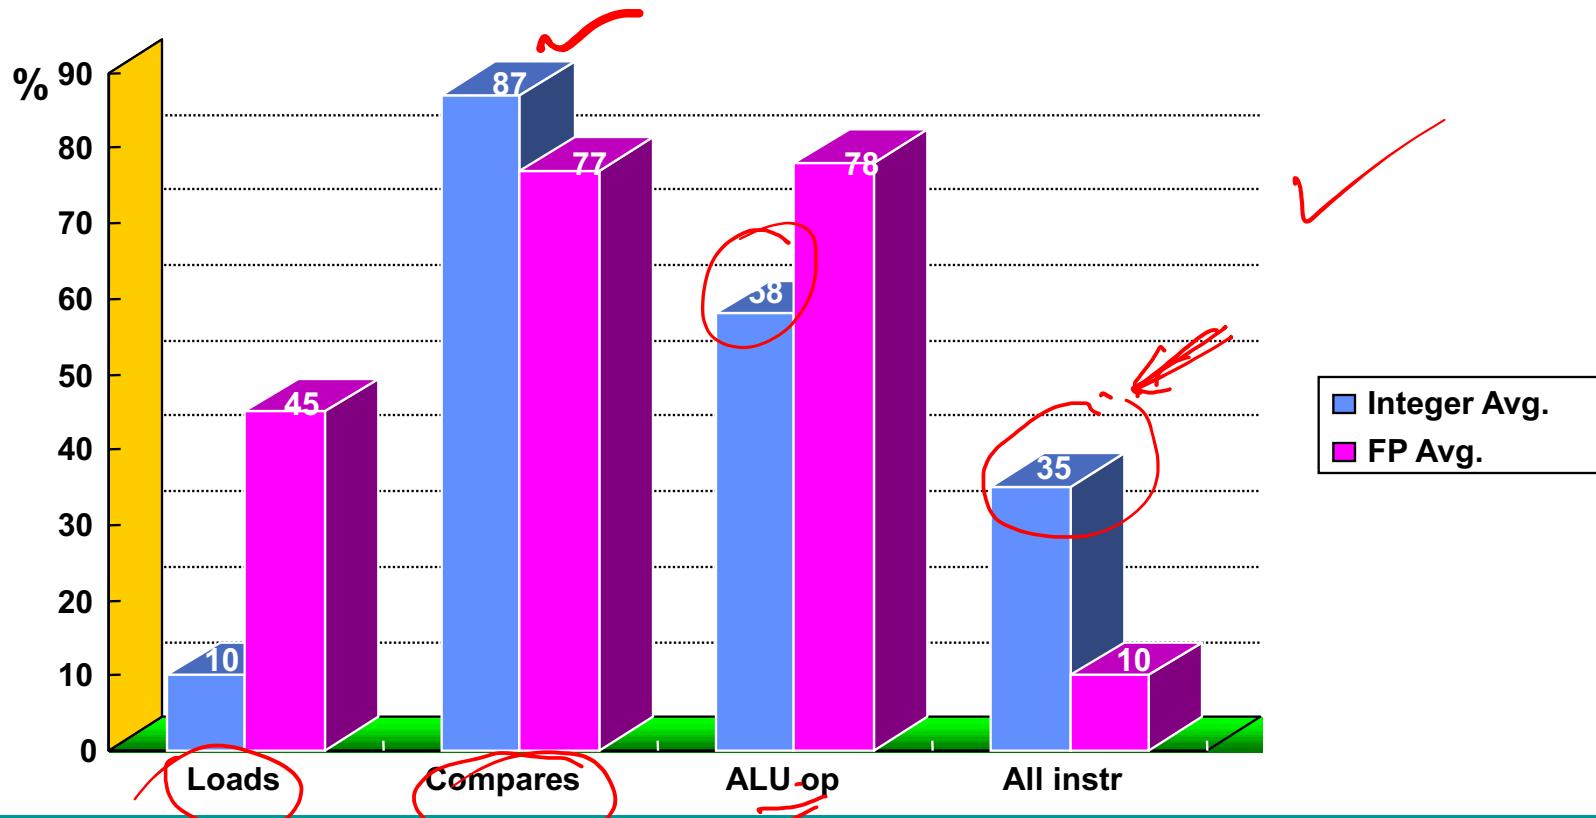
Operand Address bits: Displacement Values

- This value is related to the operand address field when the address is represented by the displacement from the base address
 - Wide distribution
 - The vast majority --- positive
 - A majority of the large displacements - negative
- \overrightarrow{A}
 \downarrow
 $A(S)$
- $=$
 \downarrow
 lwp



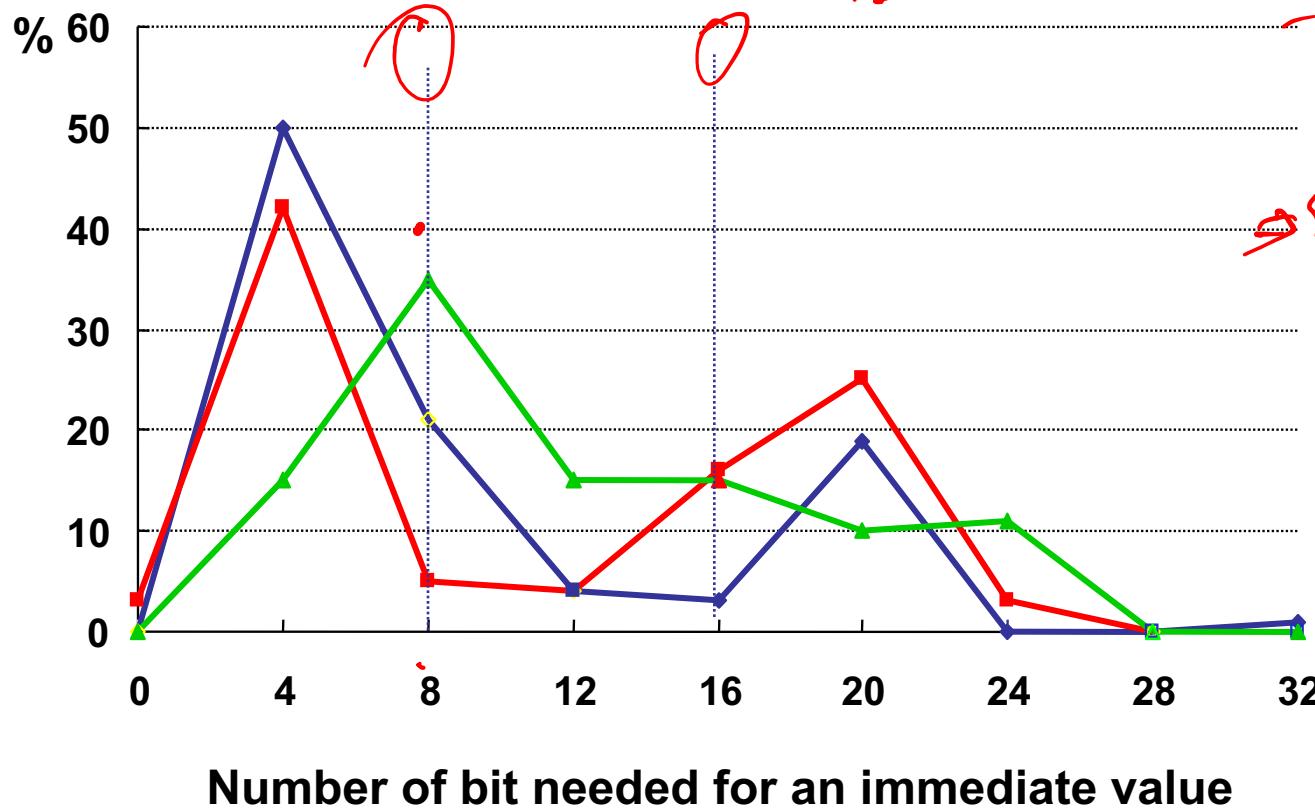
Operand Address bits: Immediate Addressing Mode

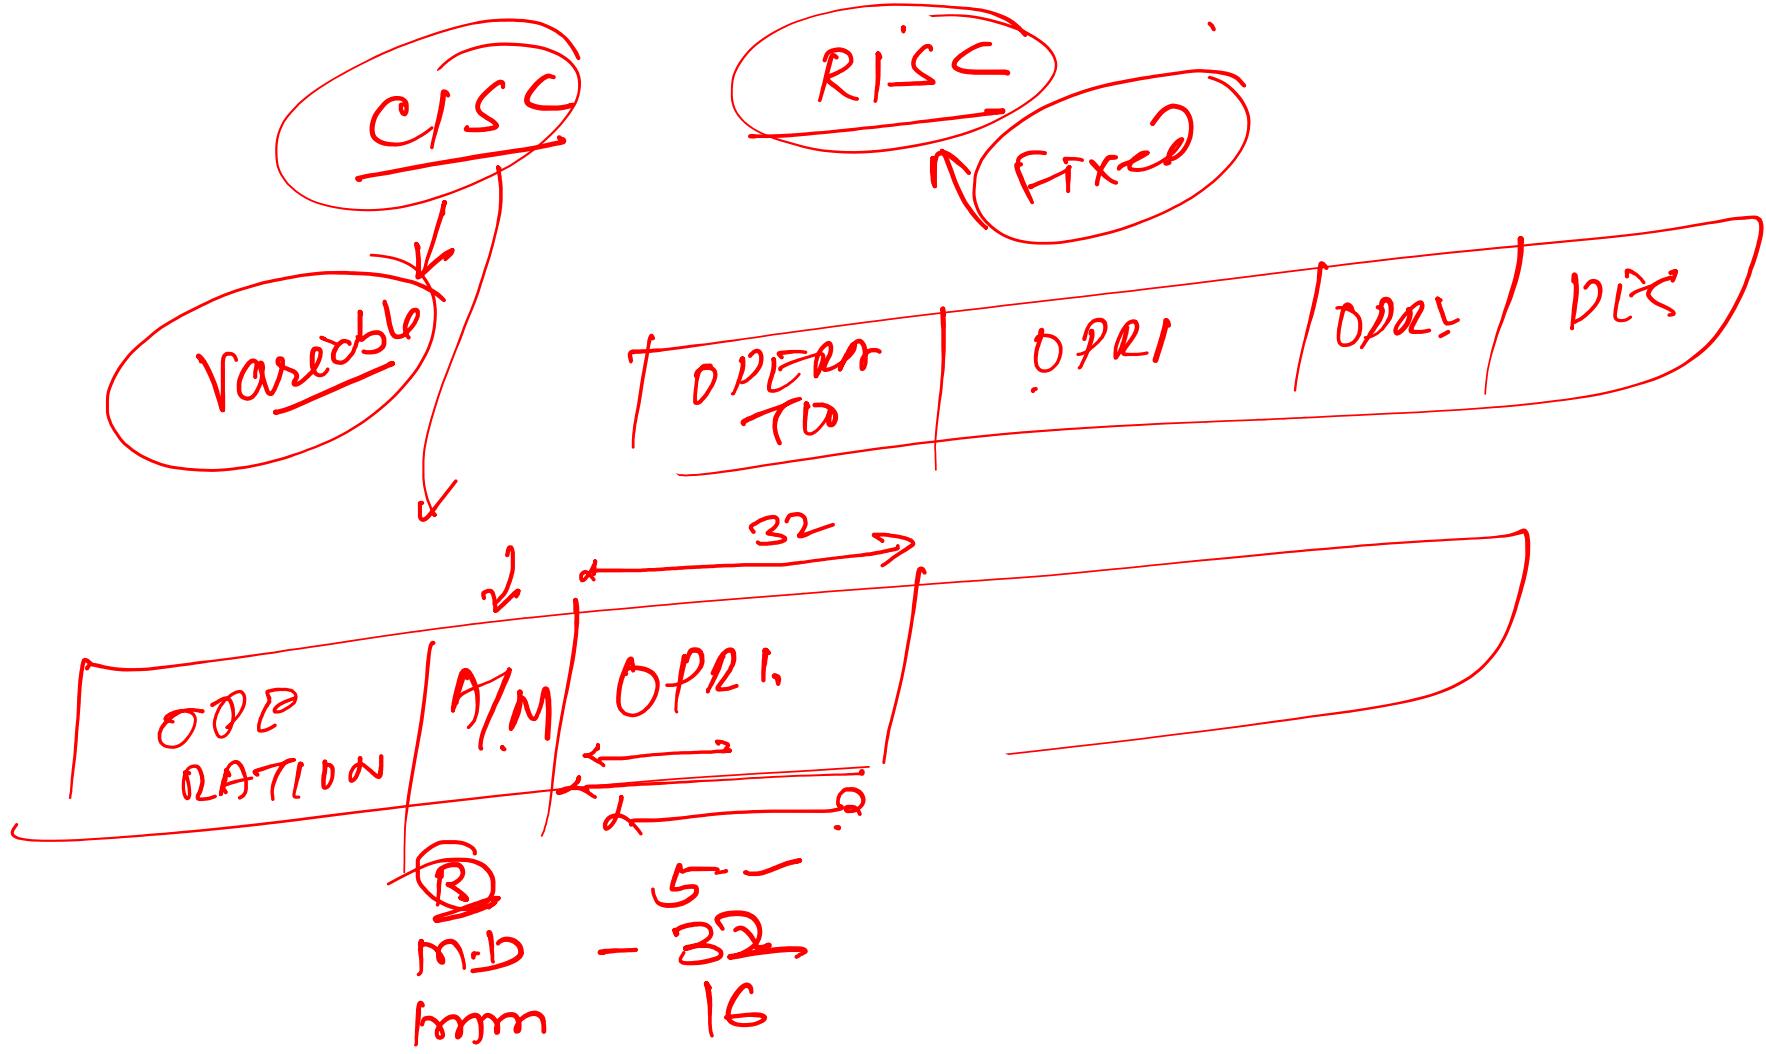
Percentage of operations that use immediates

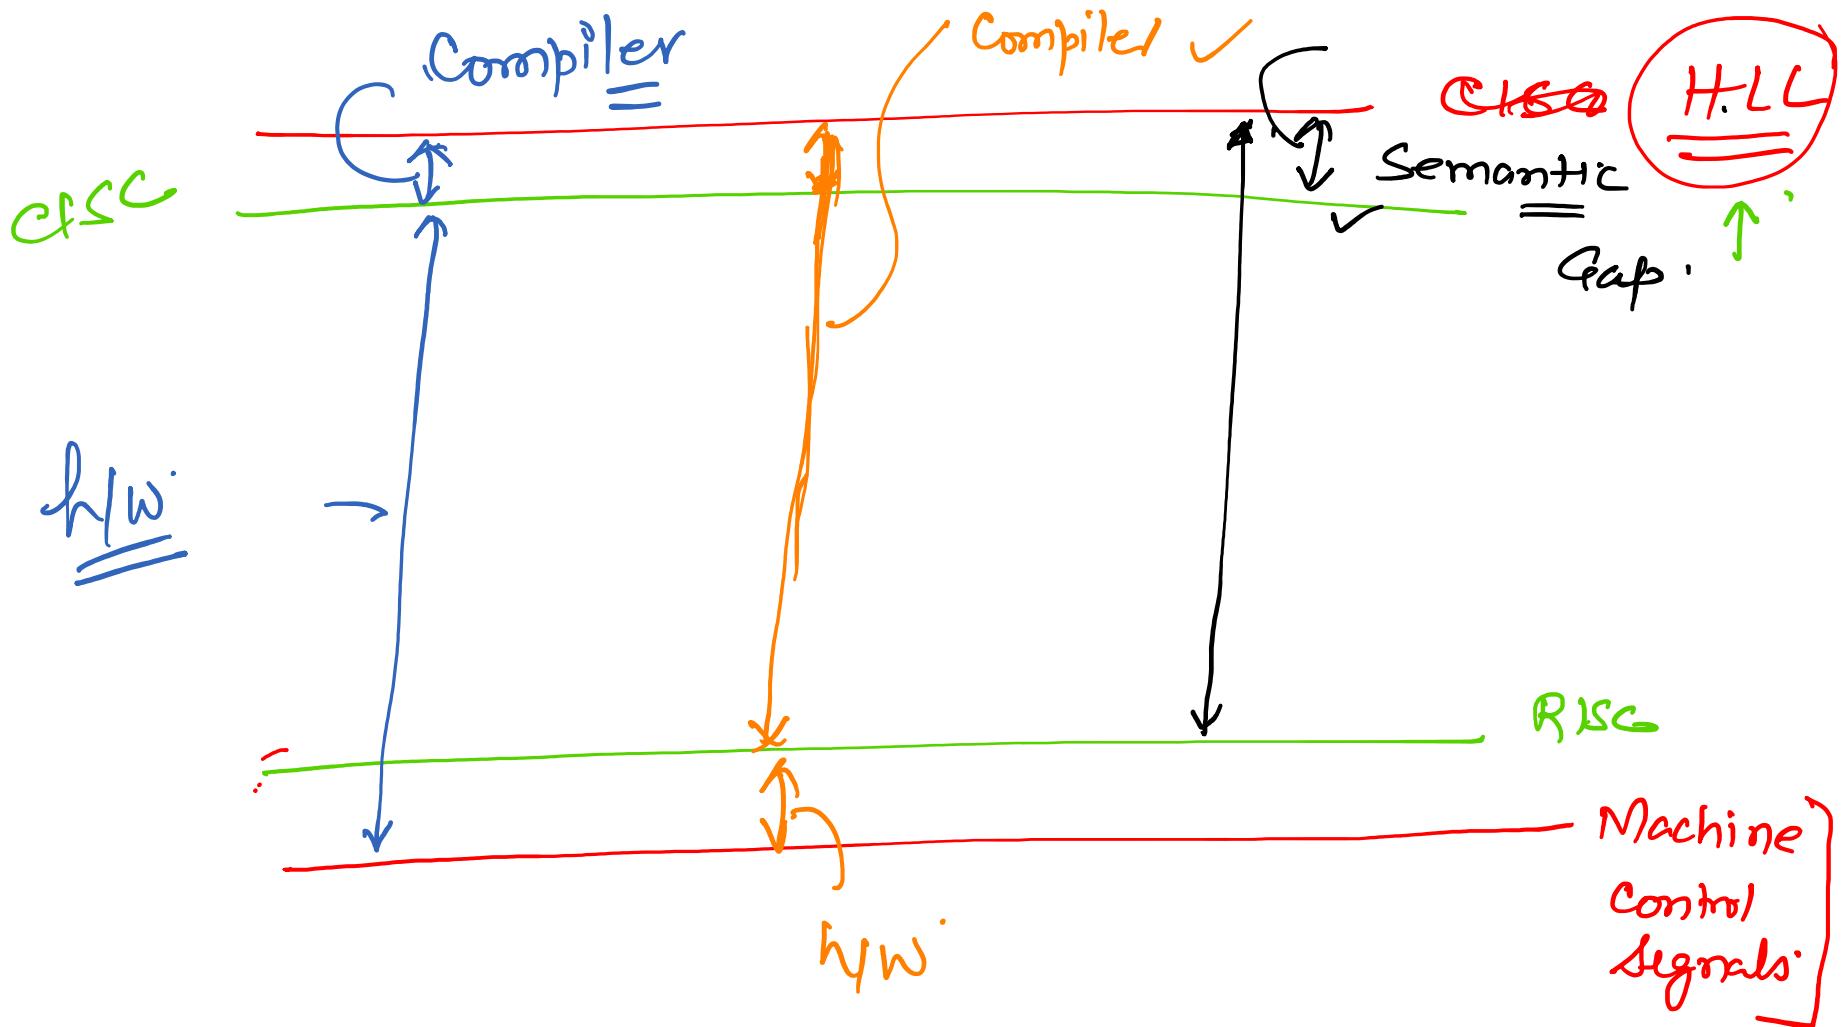


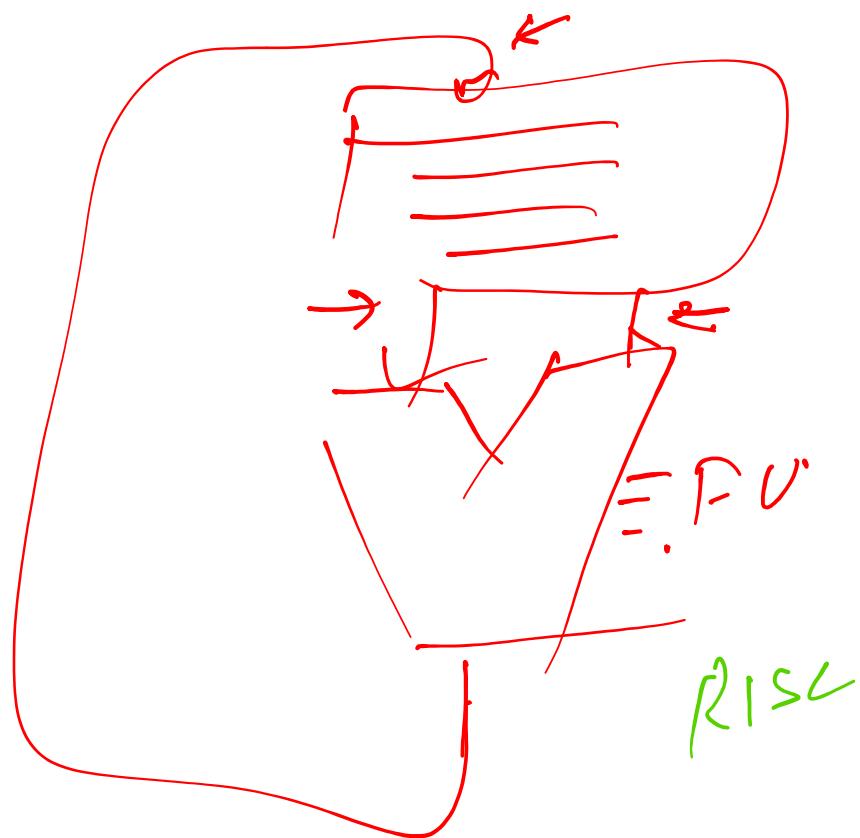
Operand Address bits: Immediate Addressing Mode

A/M
a/m
a/m

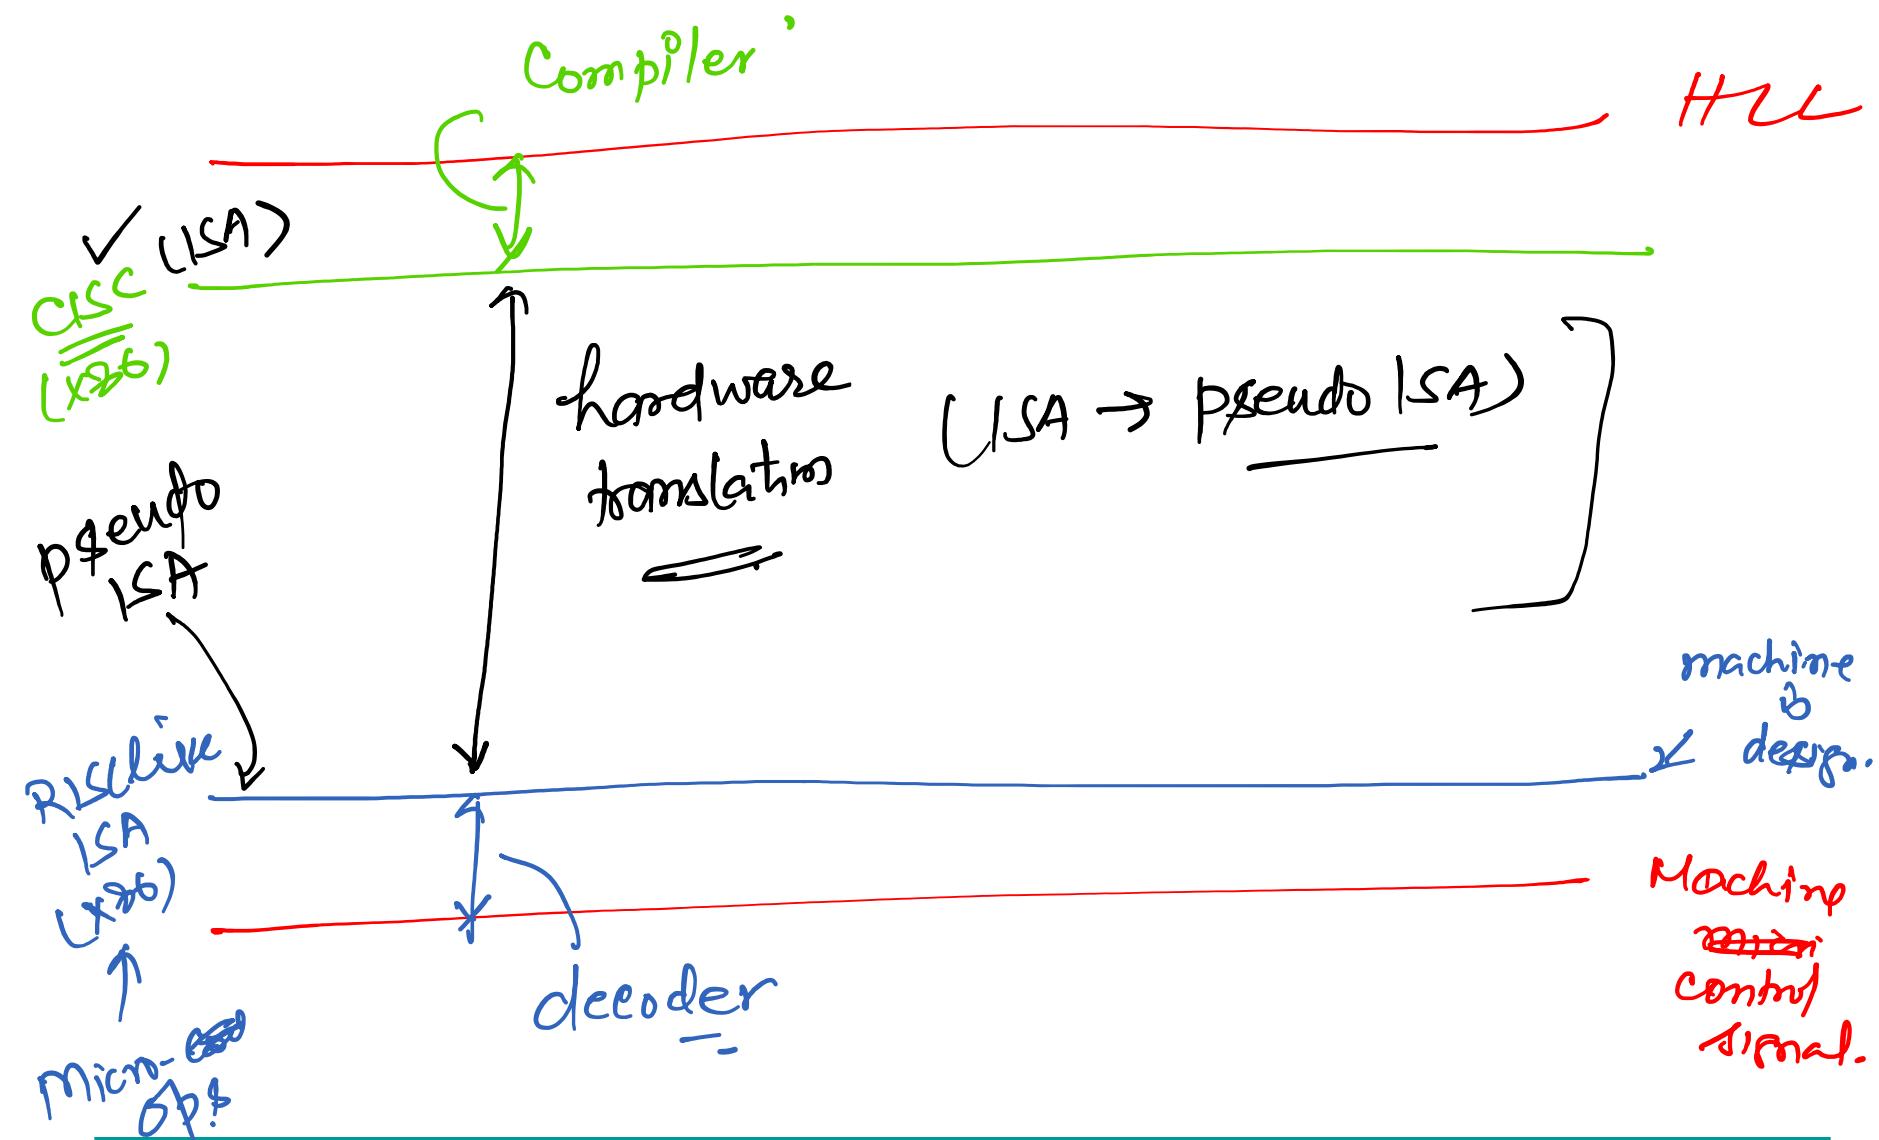








ISA Indirectios



ISA-level Tradeoffs: Semantic Gap

- Where to place the ISA? Semantic gap
 - Closer to high-level language (HLL) → Small semantic gap, complex instructions
 - Closer to hardware control signals? → Large semantic gap, simple instructions
- RISC vs. CISC machines
 - RISC: Reduced instruction set computer
 - CISC: Complex instruction set computer
 - FFT, QUICKSORT, POLY, FP instructions?
 - VAX INDEX instruction (array access with bounds checking)



ISA-level Tradeoffs: Semantic Gap

- Simple compiler, complex hardware vs. complex compiler, simple hardware
 - Caveat: Translation (indirection) can change the tradeoff!
- Burden of backward compatibility
- Performance?
 - Optimization opportunity: Example of VAX INDEX instruction: who (compiler vs. hardware) puts more effort into optimization?
 - Instruction size, code size



Small versus Large Semantic Gap

- CISC vs. RISC
 - Complex instruction set computer → complex instructions
 - Initially motivated by “not good enough” code generation
 - Reduced instruction set computer → simple instructions
 - John Cocke, mid 1970s, IBM 801
 - Goal: enable better compiler control and optimization
- RISC motivated by
 - Memory stalls (no work done in a complex instruction when there is a *memory stall*?)
 - Simplifying the hardware → lower cost, higher frequency
 - Enabling the compiler to optimize the code better
 - Find fine-grained parallelism to reduce *stalls*



How High or Low Can You Go?

- Very large semantic gap
 - Each instruction specifies the complete set of control signals in the machine
 - Compiler generates control signals
 - Open microcode (John Cocke, 1970s)
 - Gave way to optimizing compilers
- Very small semantic gap
 - ISA is (almost) the same as high-level language
 - Java machines, LISP machines, object-oriented machines, capability-based machines



Memory Address

- Interpreting memory address
 - Big Endian)
 - Little Endian)
- Instruction misalignment
- Addressing mode



Little or Big: Where to Start?

- Byte ordering: Where is the first byte?
- Big-endian : IBM, SPARC, Motorola
- Little-endian: Intel, DEC
- Supporting both: MIPS, PowerPC, ARM

Number 0x5678

Little-endian

00000003	5
00000002	6
00000001	7
00000000	8

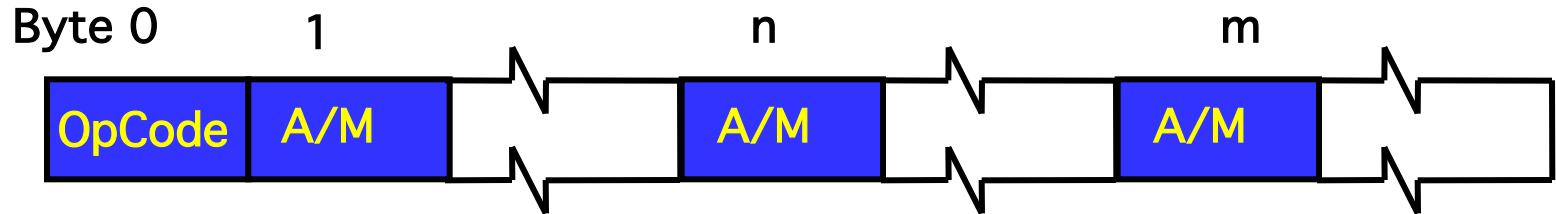
Big-endian ✓

8
7
6
5



VAX-11

Variable format, 2- and 3-address instructions



- 32-bit word size, 16 GPR (4 reserved)
- Rich set of addressing modes (apply to any operand)
- Rich set of operations
 - bit field, stack, call, case, loop, string, poly, system
- Rich set of data types (B, W, L, Q, O, F, D, G, H)
- Condition codes

ISA-level Tradeoffs: Instruction Length

- Fixed length: Length of all instructions the same
 - + Easier to decode single instruction in hardware
 - + Easier to decode multiple instructions concurrently
 - Wasted bits in instructions
 - Harder-to-extend ISA (how to add new instructions?)
- Variable length: Length of instructions different (determined by opcode and sub-opcode)
 - + Compact encoding
 - Intel 432: Huffman encoding (sort of). 6 to 321 bit instructions. **How?**
 - More logic to decode a single instruction
 - Harder to decode multiple instructions concurrently
- Tradeoffs
 - Code size (memory space, bandwidth, latency) vs. hardware complexity
 - ISA extensibility and expressiveness
 - Performance? Smaller code vs. imperfect decode



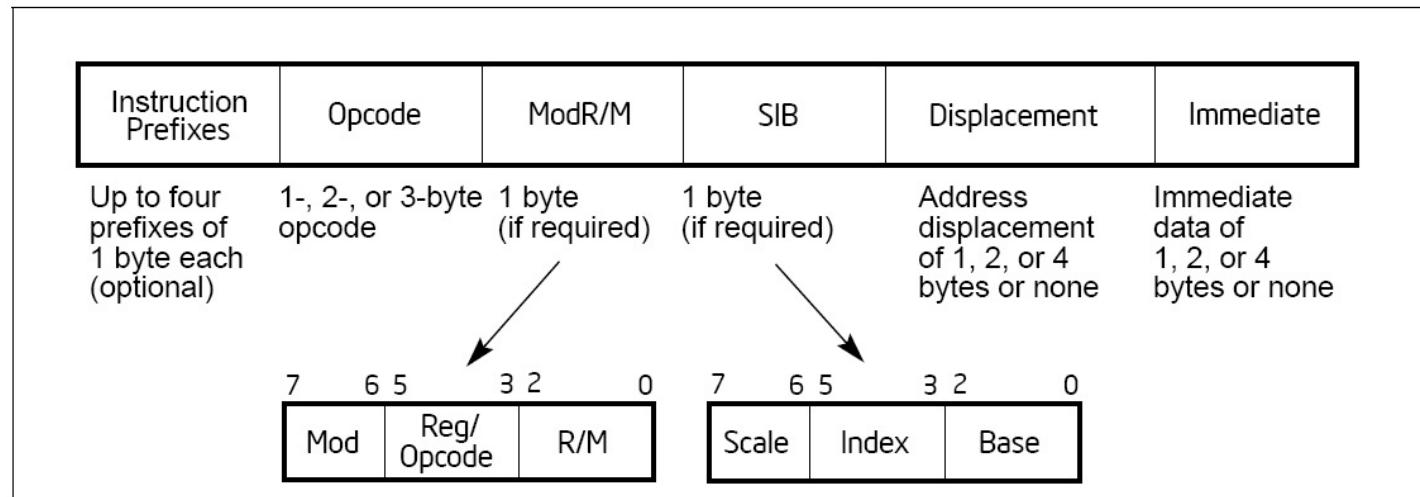
ISA-level Tradeoffs: Uniform Decode

- **Uniform decode:** Same bits in each instruction correspond to the same meaning
 - Opcode is always in the same location
 - Ditto operand specifiers, immediate values, ...
 - Many “RISC” ISAs: Alpha, MIPS, SPARC
 - + Easier decode, simpler hardware
 - + Enables parallelism: generate target address before knowing the instruction is a branch
 - Restricts instruction format (fewer instructions?) or wastes space
- **Non-uniform decode**
 - E.g., opcode can be the 1st-7th byte in x86
 - + More compact and powerful instruction format
 - More complex decode logic



x86 vs. Alpha Instruction Formats

- x86:



- Alpha:

31	26 25	21 20	16 15	5 4	0	PALcode Format
Opcode			Number			Branch Format
Opcode	RA		Disp			Memory Format
Opcode	RA	RB	Disp			Operate Format
Opcode	RA	RB	Function	RC		



MIPS Instruction Format

- R-type, 3 register operands

0	rs	rt	rd	shamt	funct
6-bit	5-bit	5-bit	5-bit	5-bit	6-bit

R-type

- I-type, 2 register operands and 16-bit immediate operand

opcode	rs	rt	immediate
6-bit	5-bit	5-bit	16-bit

I-type

- J-type, 26-bit immediate operand

opcode	immediate
6-bit	26-bit

J-type

- Simple Decoding

- 4 bytes per instruction, regardless of format
- must be 4-byte aligned (2 lsb of PC must be 2b'00)
- format and fields easy to extract in hardware



Instruction Execution Characteristics: Type of Operations

Relative Dynamic Frequencies of statements in HLL programs

What type of statements is most frequent?

- Assignment statements dominate
 - Functional instructions and Transfer instructions
 - Movements of data must be made simple, thus fast
- Conditional Statements(if and loop together)
 - Instructions with Control function
 - Sequence control mechanism is important



Instruction Execution Characteristics: Time Consumed by Statements

Time Consumed \longleftrightarrow Number of Machine Instructions

	Dynamic Occur		Machine Instr Wt		Memory Ref Wt	
	PASCAL	C	PASCAL	C	PASCAL	C
Assignment	45	38	13	13	14	15
Loop	5	3	42	32	33	26
Call/Return	15	12	31	33	44	45
If	29	43	11	21	7	13
goto	-	3	-	-	-	-
others	6	1	3	1	2	1

Machine instruction weighted

$$= [\text{Average No. of machine Instr. / Statements}] \times [\text{Frequency of Occurrences}]$$

Memory reference weighted

$$= [\text{Average No. of memory references / Statement}] \times [\text{Frequency of Occurrences}]$$

Most time consuming statement is procedure CALL/RETURN



Instruction Execution Characteristics: Type of Operands

Dynamic Frequencies of Occurrences

	PASCAL	C	Average
Integer Constant	16	23	20
Scalar Variable	58	53	55
Array/Structure	26	24	25

Majority of references to scalar

- 80% are local to a procedure
- References to arrays/structure require index or pointer

Locations of operands(Average per instruction)

- 0.5 operands in memory
- 1.4 operands in registers

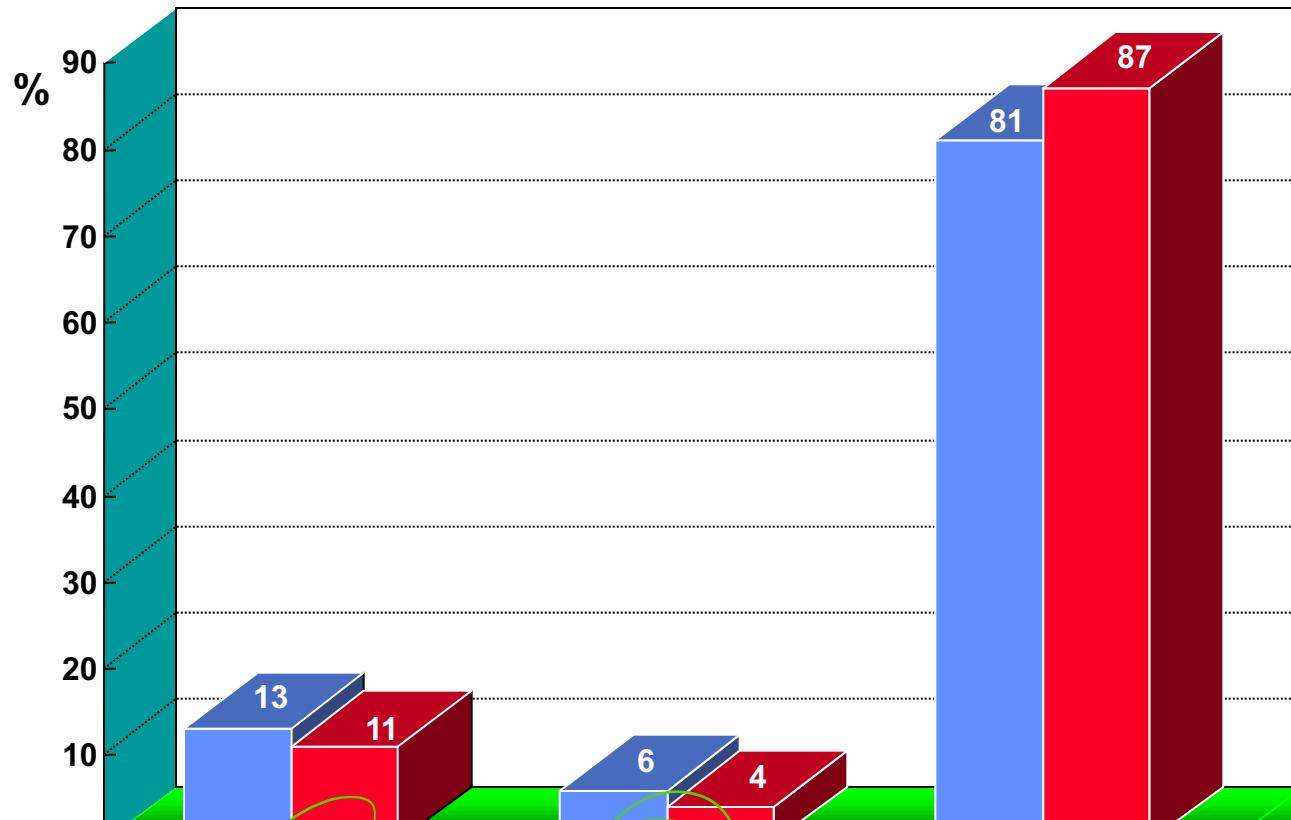


Operations in the Instr. Set

Rank	80x86 instructions	Integer average (% total executed)
1	load	22%
2	conditional branch	20%
3	compare	16%
4	store	12%
5	add	8%
6	and	6%
7	sub	5%
8	move reg-reg	4%
9	call	1%
10	return	1%
Total		96%



Control Flow Instructions



Pipelined
Arch

Thank You

