

Pipelined Architecture

Virendra Singh

Computer Architecture and Dependable Systems Lab

Department of Electrical Engineering

Indian Institute of Technology Bombay

<http://www.ee.iitb.ac.in/~viren/>

E-mail: viren@ee.iitb.ac.in

JSA
70 - 80's

RISC
CISC

CS-683: Advanced Computer Architecture



Lecture 6 (10 Aug 2021)

CADSL

RISC Architecture

- Simple instructions
- Fixed Instruction Encoding
- Limited Addressing Mode
- Instruction count increases ✓
- Simple controller ✓
- Load/Store architecture ✓
- Limited addressing modes

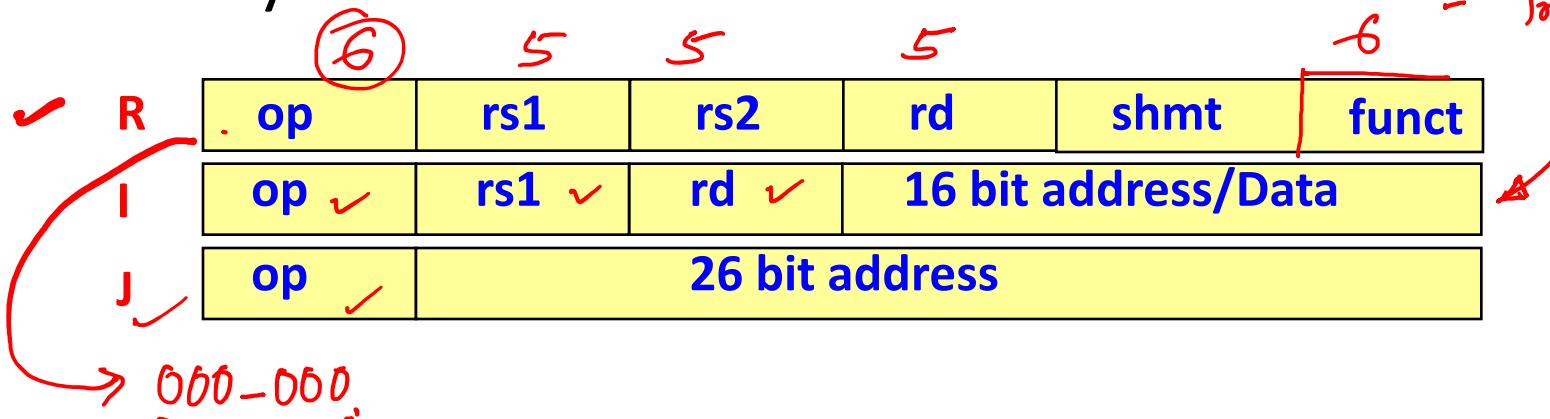


Overview of MIPS

- ❖ Simple instructions, all 32 bits wide
- ❖ Very structured, no unnecessary baggage
- ❖ Only three instruction formats

~~add
a < 100~~

- Reg Add.
- Immediate



$$2^6 = \cancel{32} \ 64$$



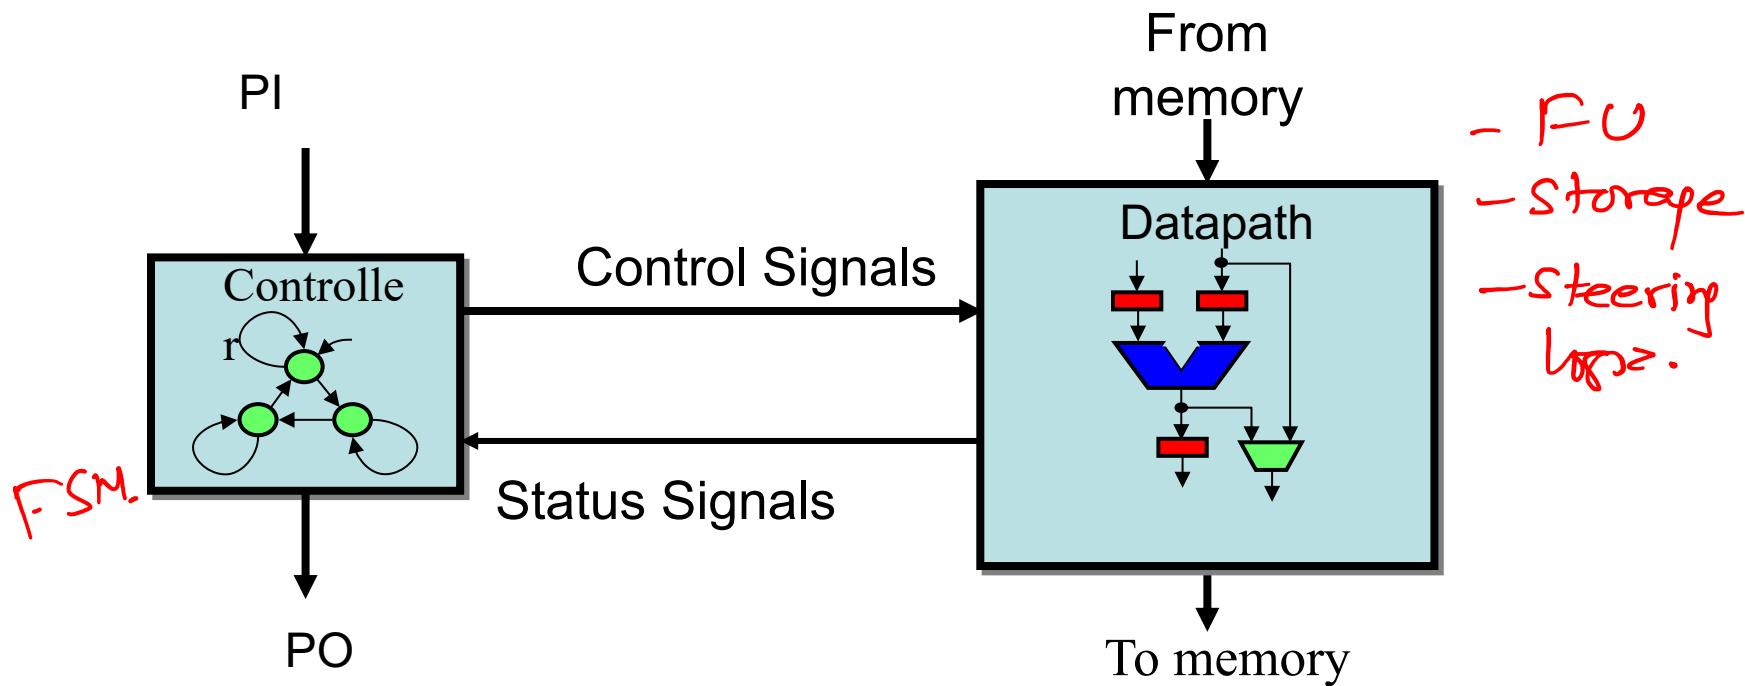
Example processor MIPS subset

MIPS Instruction – Subset

- ❖ Arithmetic and Logical Instructions
 - add, sub, or, and, slt ✓
- ❖ Memory reference Instructions
 - lw, sw ✓
- ❖ Branch
 - beq, j ✓



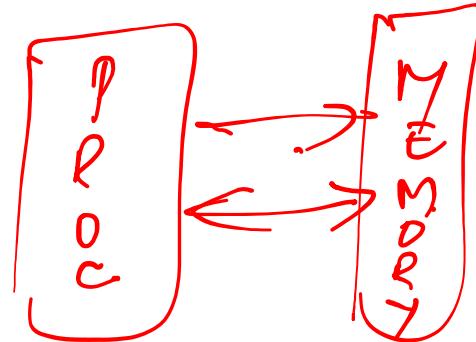
Processor Architecture



Simple

ADD / SUB / MUL

- ① fetch an instruction
(PC / IP)
- ② Understand (Decode)
- ③ Read Operands
- ④ Execute instruction
- ⑤ Update the result



Can complete
in one step

add. $r_1, r_2 \rightarrow r_3$

$$T = t_{inc} = t_f + t_d + t_{or} + t_e + t_u$$

$$r_1 = \frac{r_2 + r_3}{r}$$

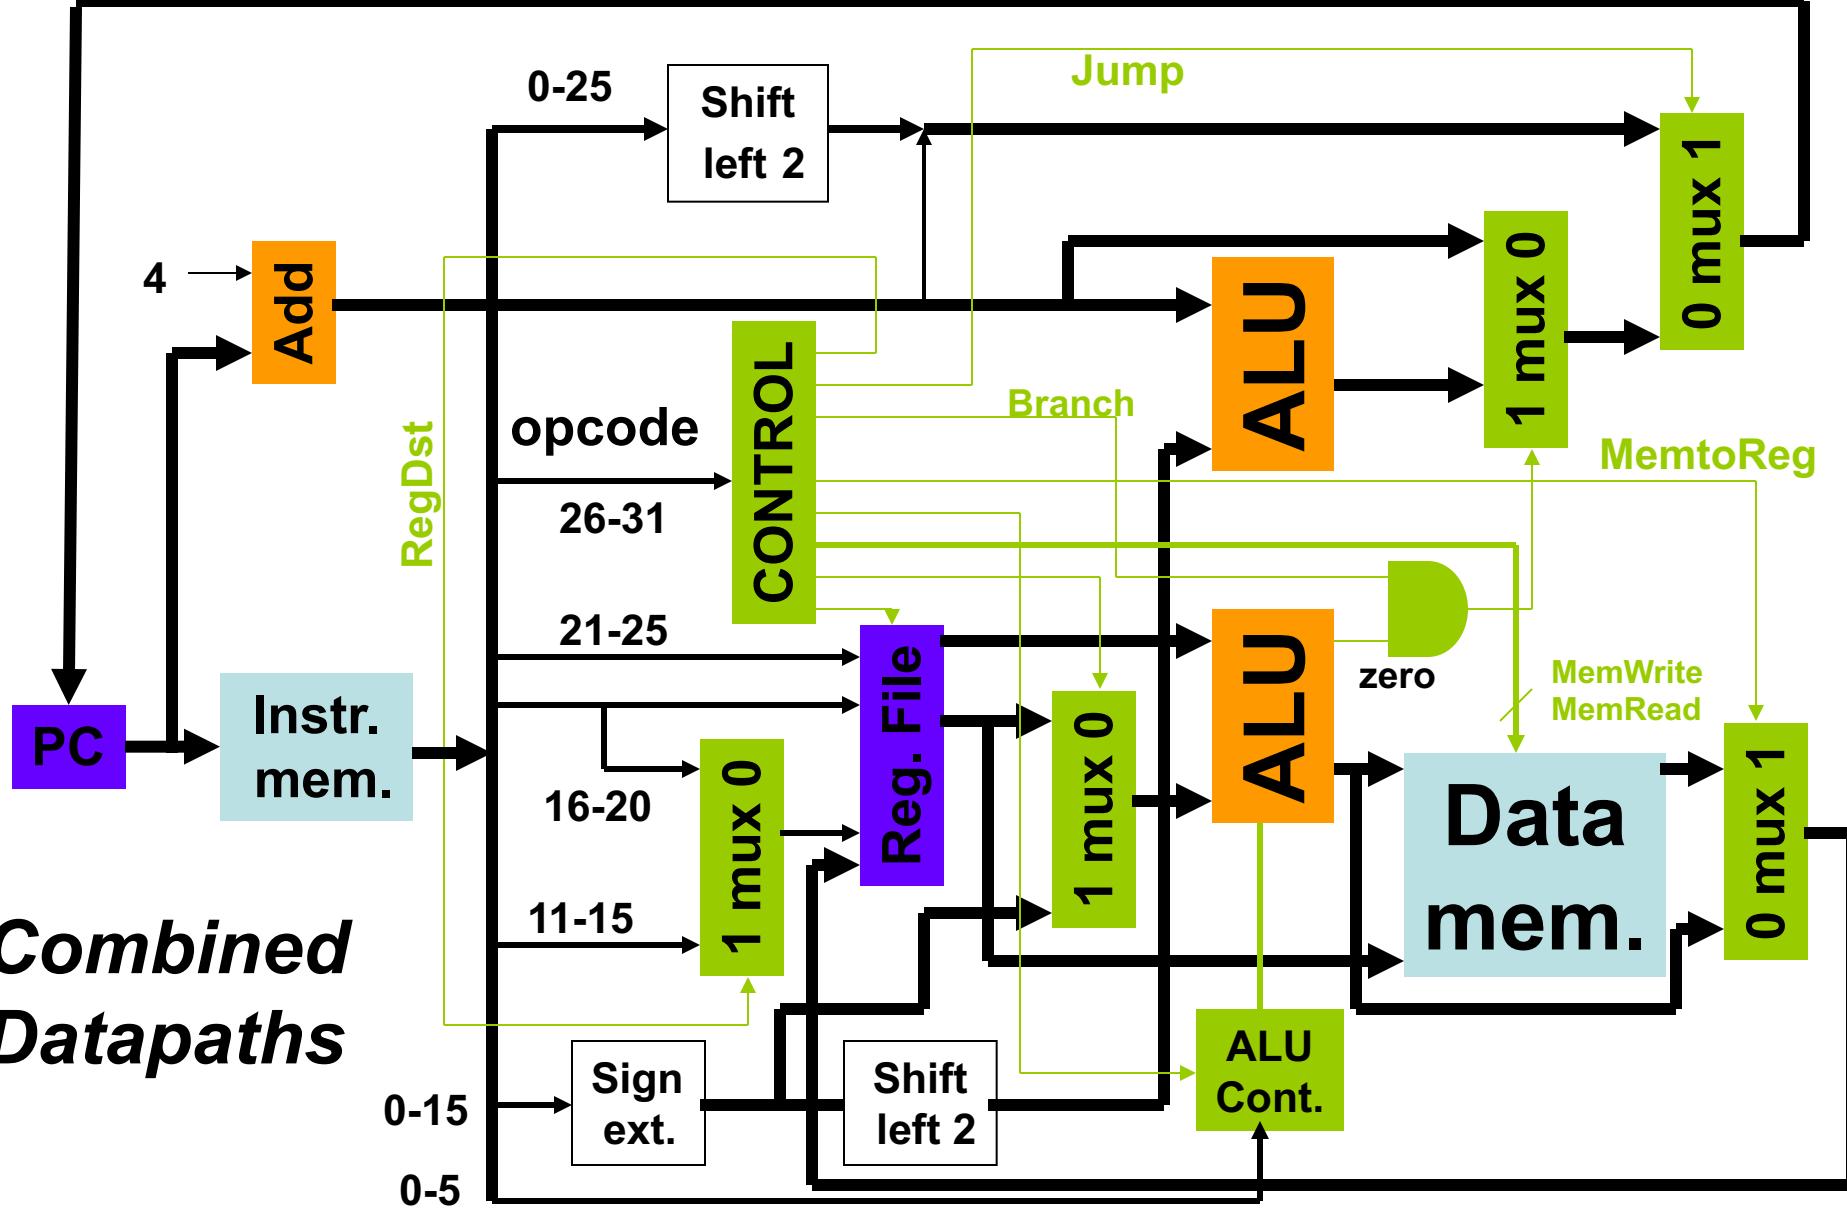
Single cycle implementation CPI=1



Time , $T = \underline{\underline{C}} \times \overbrace{CPI \times \underline{\underline{C}}}$



Combined Datapaths



Time for Jump (J-Type)

- ALU (R-type) 6ns
- Load word (I-type) 8ns
- Store word (I-type) 7ns
- Branch on equal (I-type) 5ns
- Jump (J-type)
 - Fetch (memory read) 2ns
 - Total 2ns



- 1. fetch
- 2. understand (decode)
- 3. operand read
- 4. Execute ↗
- 5. update the result

One step → only one sub task.

Multiple steps ↗

$$\text{Time} = \max\{t_f + t_d, t_{or}, t_e, t_u\}$$

$$\underline{\underline{CPI}} > 1$$

Mult-cycle Implementation



AL

1. fetch 2. decode 3. Op. read. 4. Ex- 5. Update- } 5

Load

1. fetch, 2. Decode 3. Op. read. 4. Add computation, } 6
5. Read memmey 6. Update }

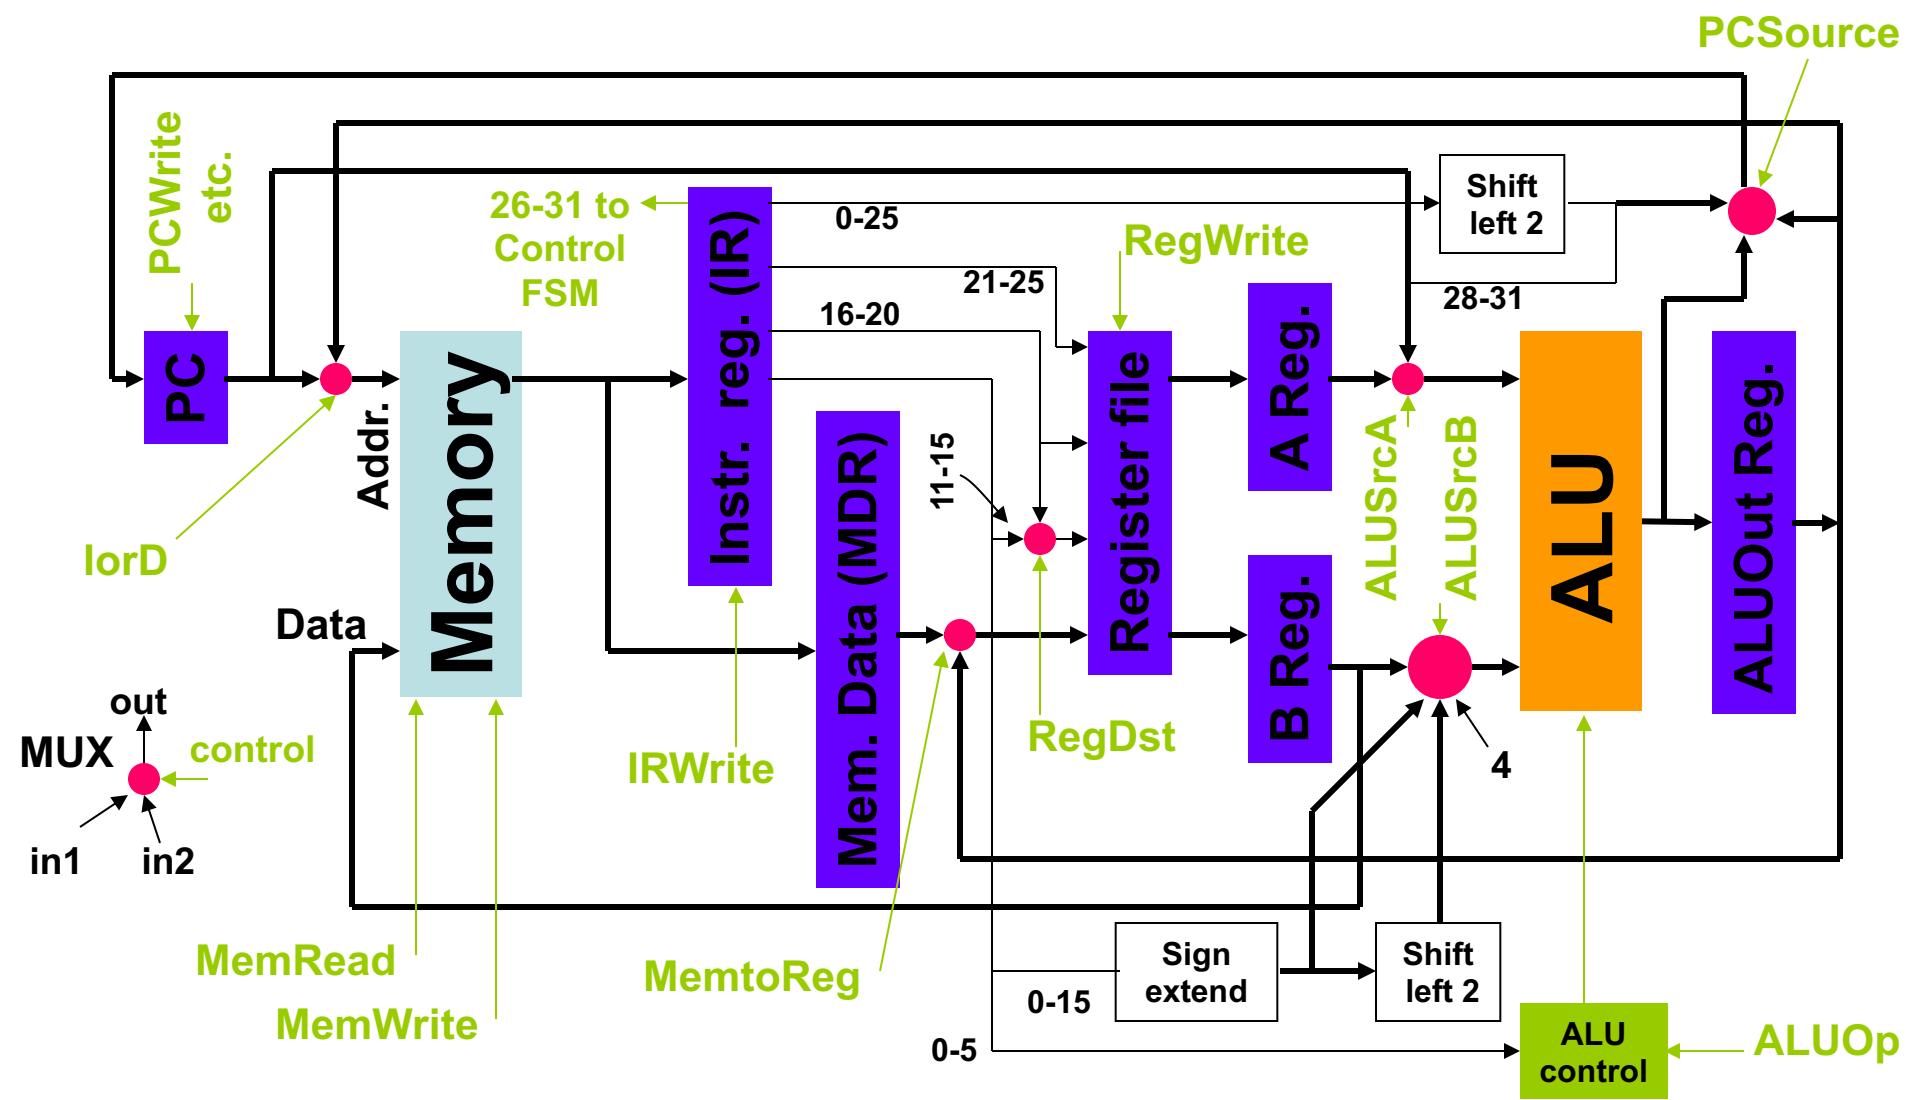


Multicycle Instruction Execution

Step	R-type (4 cycles)	Mem. Ref. (4 or 5 cycles)	Branch type (3 cycles)	J-type (3 cycles)
Instruction fetch	IR \leftarrow Memory[PC]; PC \leftarrow PC+4			
Instr. decode/ Reg. fetch	A \leftarrow Reg(IR[21-25]); B \leftarrow Reg(IR[16-20]) ALUOut \leftarrow PC + (sign extend IR[0-15]) << 2			
Execution, addr. Comp., branch & jump completion	ALUOut \leftarrow A op B	ALUOut \leftarrow A+sign extend (IR[0-15])	If (A = B) then PC \leftarrow ALUOut	PC \leftarrow PC[28- 31] (IR[0-25]<<2)
Mem. Access or R-type completion	Reg(IR[11- 15]) \leftarrow ALUOut	MDR \leftarrow M[ALUout] or M[ALUOut] \leftarrow B		
Memory read completion		Reg(IR[16-20]) \leftarrow MDR		



Multicycle Datapath



CPI of a Computer

$$\text{CPI} = \frac{\sum_k (\text{Instructions of type } k) \times \text{CPI}_k}{\sum_k (\text{instructions of type } k)}$$

where

CPI_k = Cycles for instruction of type k

Note: *CPI is dependent on the instruction mix of the program being run. Standard benchmark programs are used for specifying the performance of CPUs.*



50% ALU
CPI = 5
- 5x5 +
20% Load.
• 2x6.

What Next

Simple cycle.

$$\text{CPI} = 1$$

$T \uparrow$

Multicycle

$$\text{CPI} \uparrow$$

$C \downarrow$

$$\text{CPI} = 1$$



Traffic Flow



Single Lane Traffic



© BNPS.CO.UK



10 Aug 2021

CS683@IITB

17

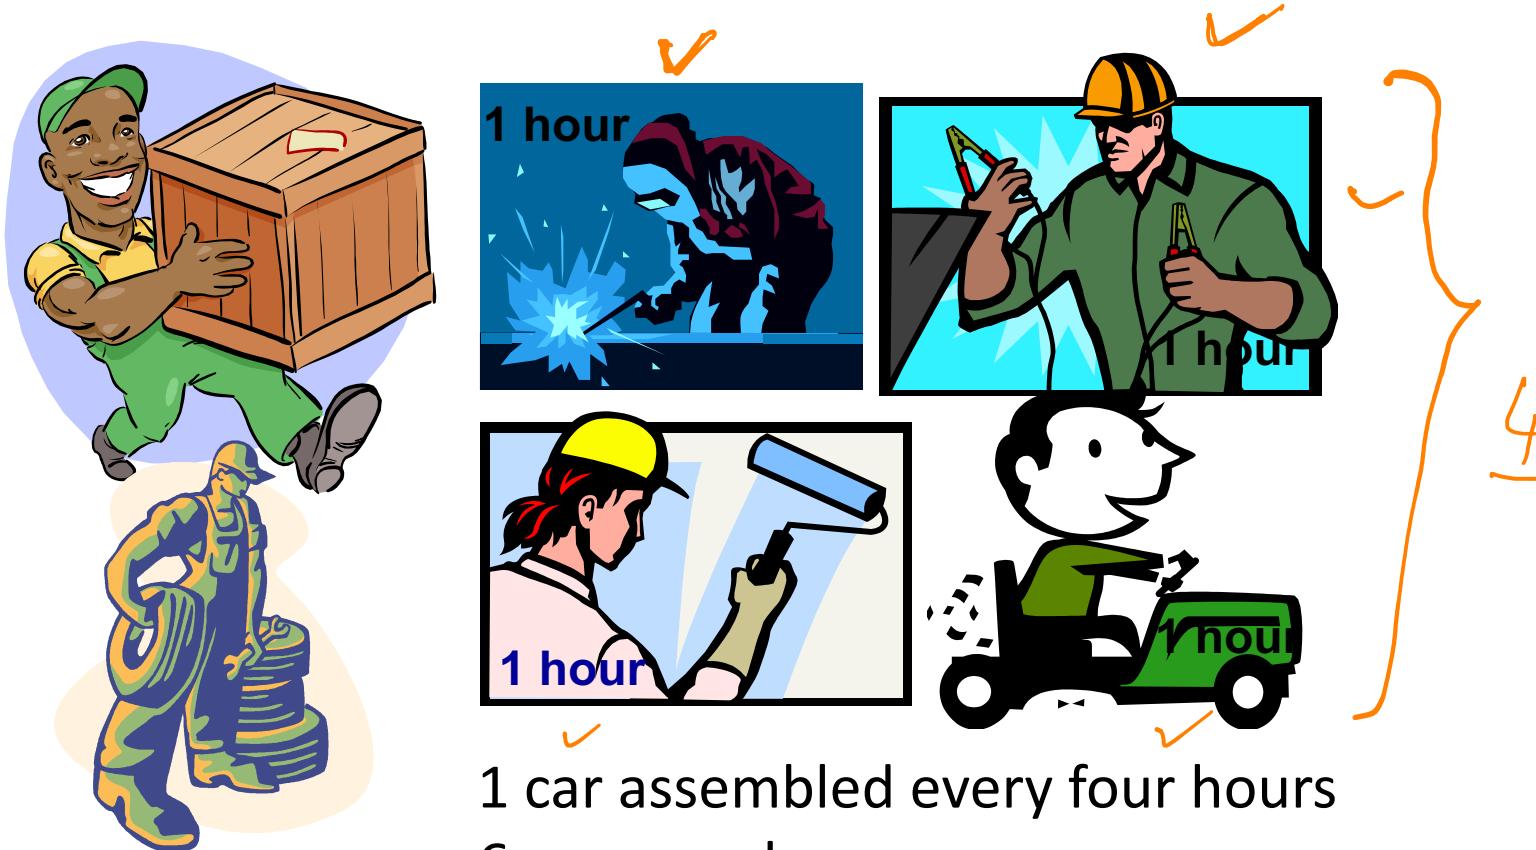
CADSL

ILP: Instruction Level Parallelism

- Single-cycle and multi-cycle datapaths execute one instruction at a time.
- How can we get better performance?
- Answer: Chaining of instructions
- Takes advantage of
 - CPI of single cycle ✓
 - Cycle time of multicycle implementation

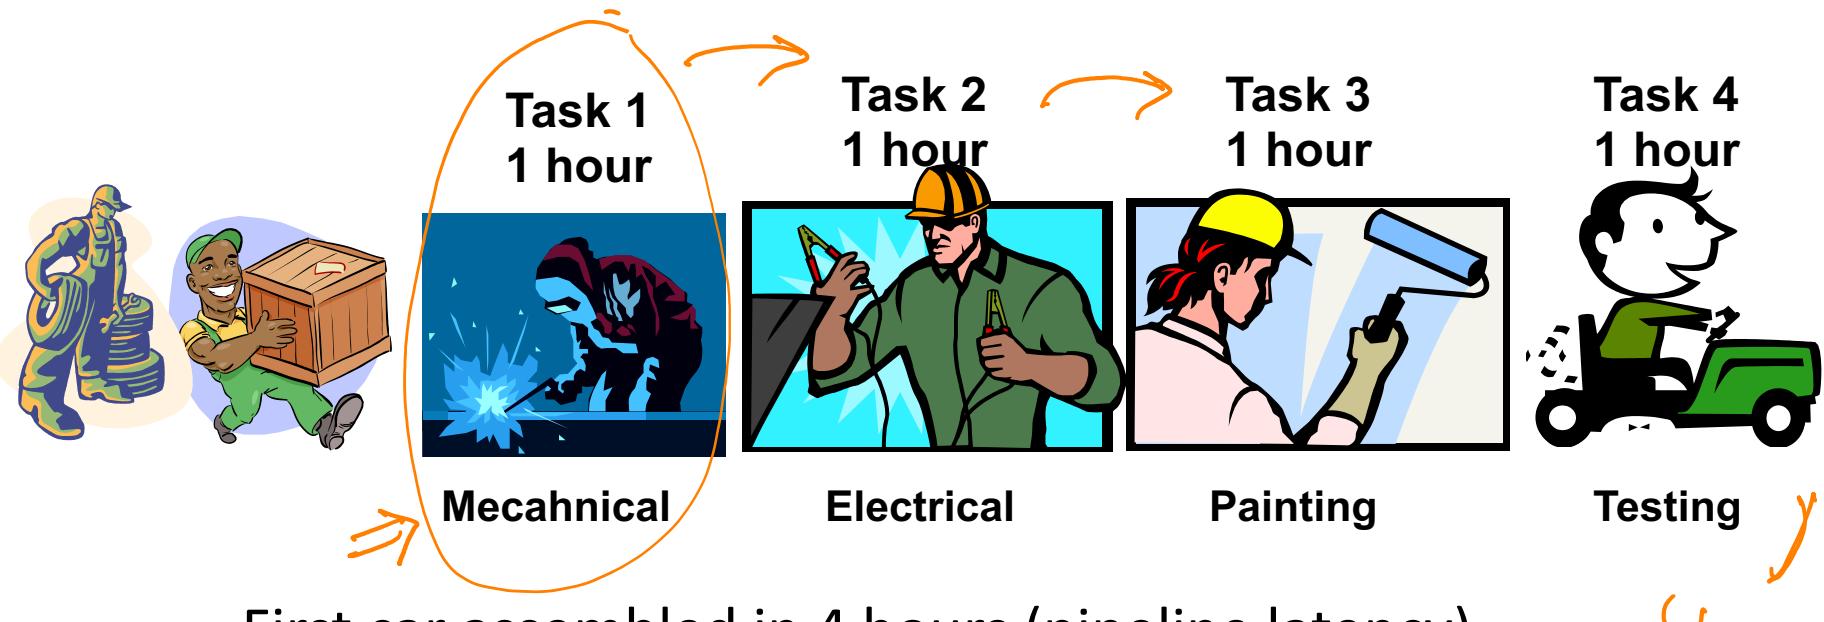


Automobile Team Assembly



1 car assembled every four hours
6 cars per day
180 cars per month
2,040 cars per year

Automobile Assembly Line



First car assembled in 4 hours (pipeline latency)

thereafter 1 car per hour

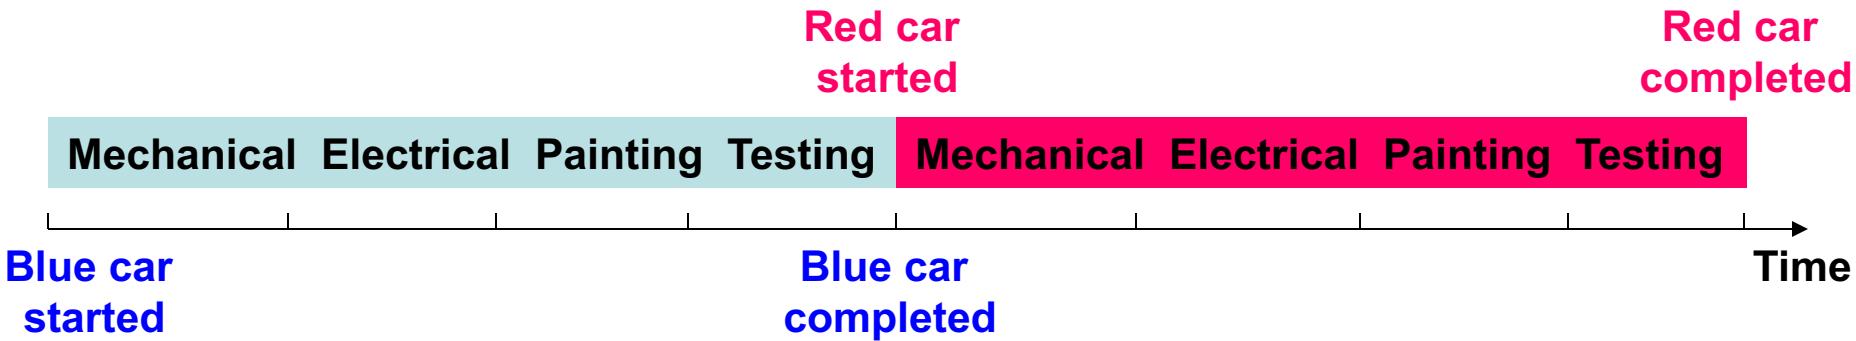
21 cars on first day, thereafter 24 cars per day

717 cars per month ✓

8,637 cars per year ✓



Throughput: Team Assembly



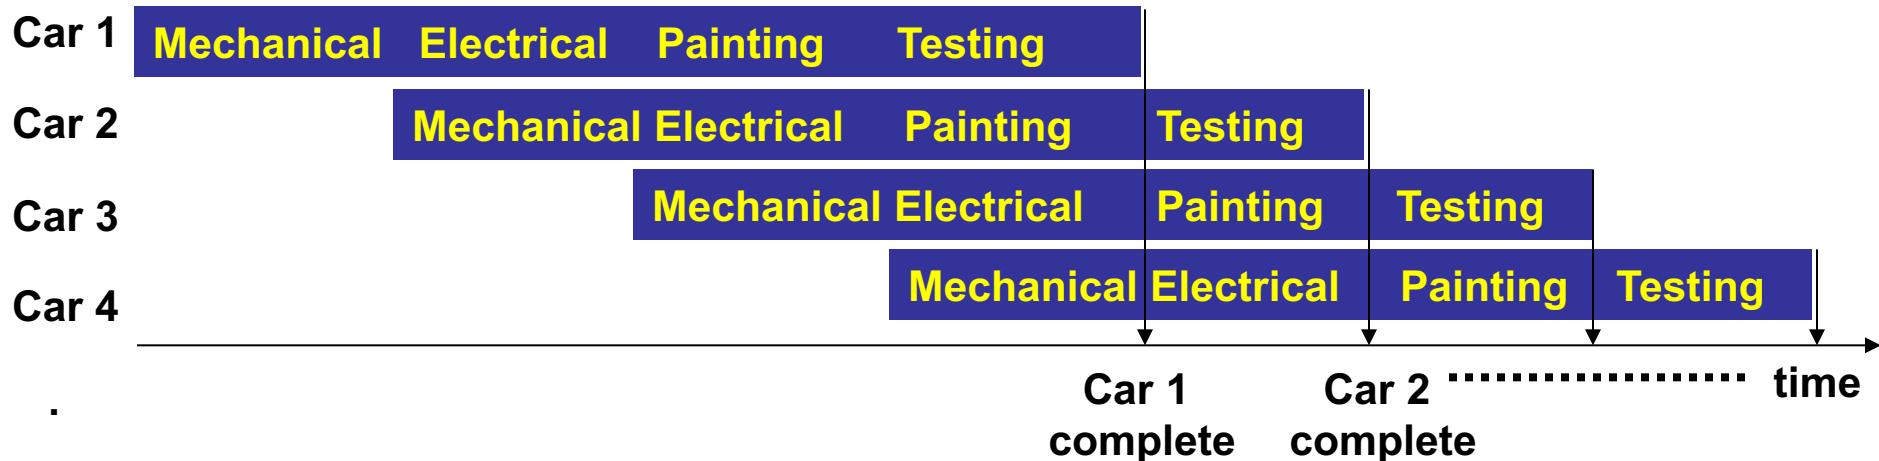
Time of assembling one car = n hours

where n is the number of nearly equal subtasks,
each requiring 1 unit of time

Throughput = $\frac{1}{n}$ cars per unit time



Throughput: Assembly Line



Time to complete first car = n time units (latency)

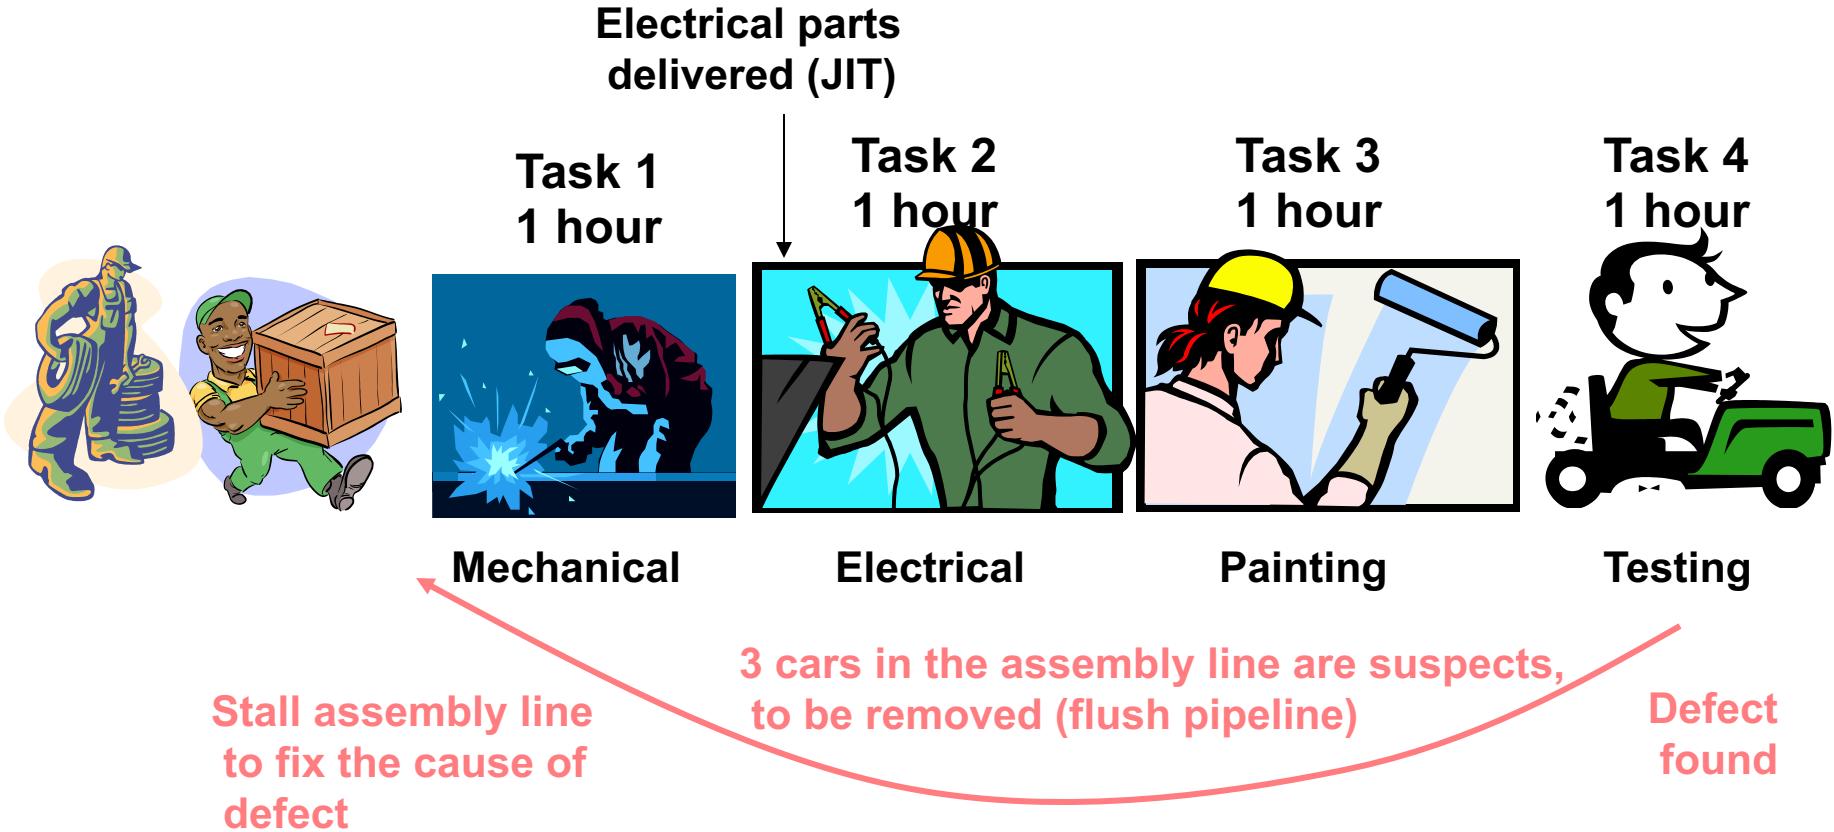
Cars completed in time T = $T - n + 1$

Throughput = $1 - (n - 1)/ T$ car per unit time

$$\frac{\text{Throughput (assembly line)}}{\text{Throughput (team assembly)}} = \frac{1 - (n - 1)/ T}{1/n} = n - \frac{n(n - 1)}{T} \rightarrow n \quad \text{as } T \rightarrow \infty$$



Some Features of Assembly Line



Pipelining in a Computer

- Divide datapath into nearly equal tasks, to be performed serially and requiring non-overlapping resources.
- Insert registers at task boundaries in the datapath; registers pass the output data from one task as input data to the next task.
- Synchronize tasks with a clock having a cycle time that just exceeds the time required by the longest task.
- Break each instruction down into a fixed number of tasks so that instructions can be executed in a staggered fashion.



AL

fetch
Decode &
Operando
read
Execute
Update

(load)

fetch
Decode &
OR
Execute
(add.
computatio)
Mem.
Access
(read)
Update

Store

fetch
Decode
Execute
Memory
Write



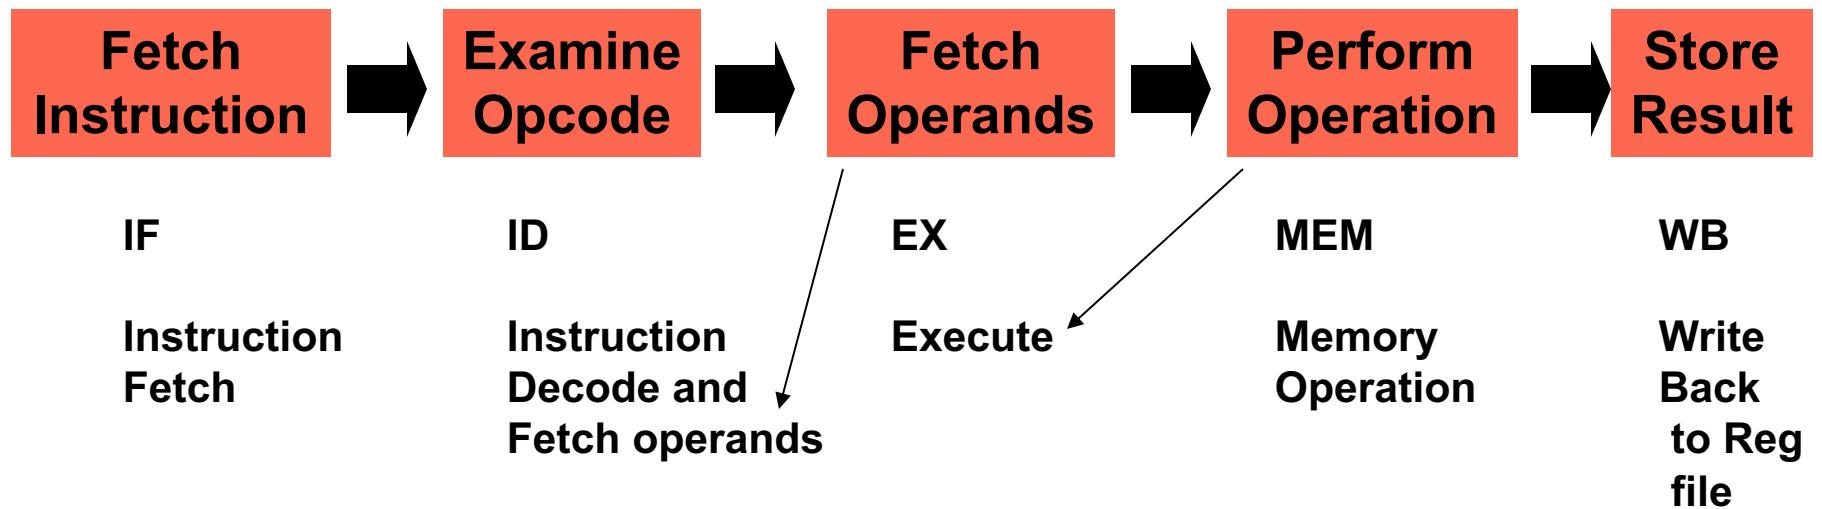
Pipelined Datapath

Instruction class	Instr. fetch (IF)	Instr. Decode (also reg. file read) (ID)	Execution (ALU Operation) (EX)	Data access (MEM)	Write Back (Reg. file write) (WB)	Total time
lw	2ns	1ns 2ns	2ns	2ns	1ns 2ns	10ns
sw	2ns	1ns 2ns	2ns	2ns	1ns 2ns	10ns
R-format: add, sub, and, or, slt	2ns	1ns 2ns	2ns	2ns	1ns 2ns	10ns
B-format: beq	2ns	1ns 2ns	2ns	2ns	1ns 2ns	10ns

No operation on data; idle time inserted to equalize instruction lengths.



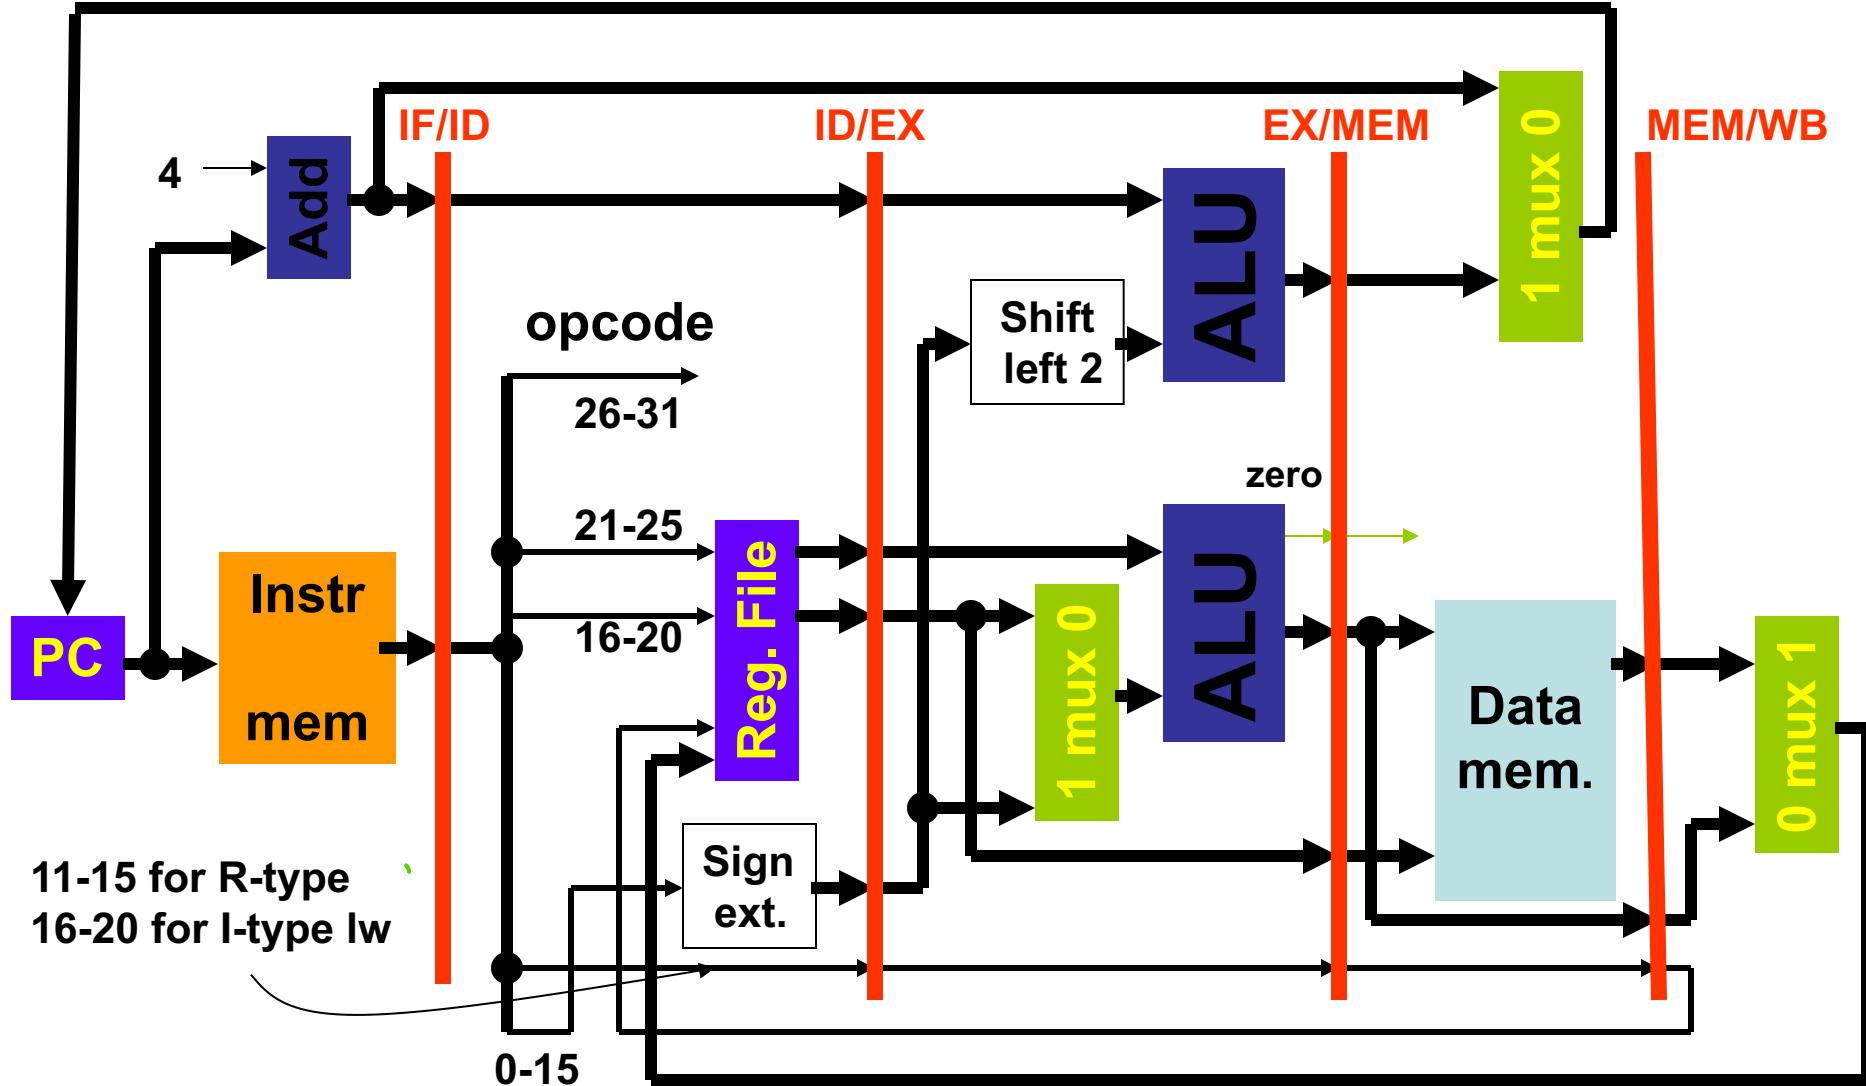
Pipelining of RISC Instructions



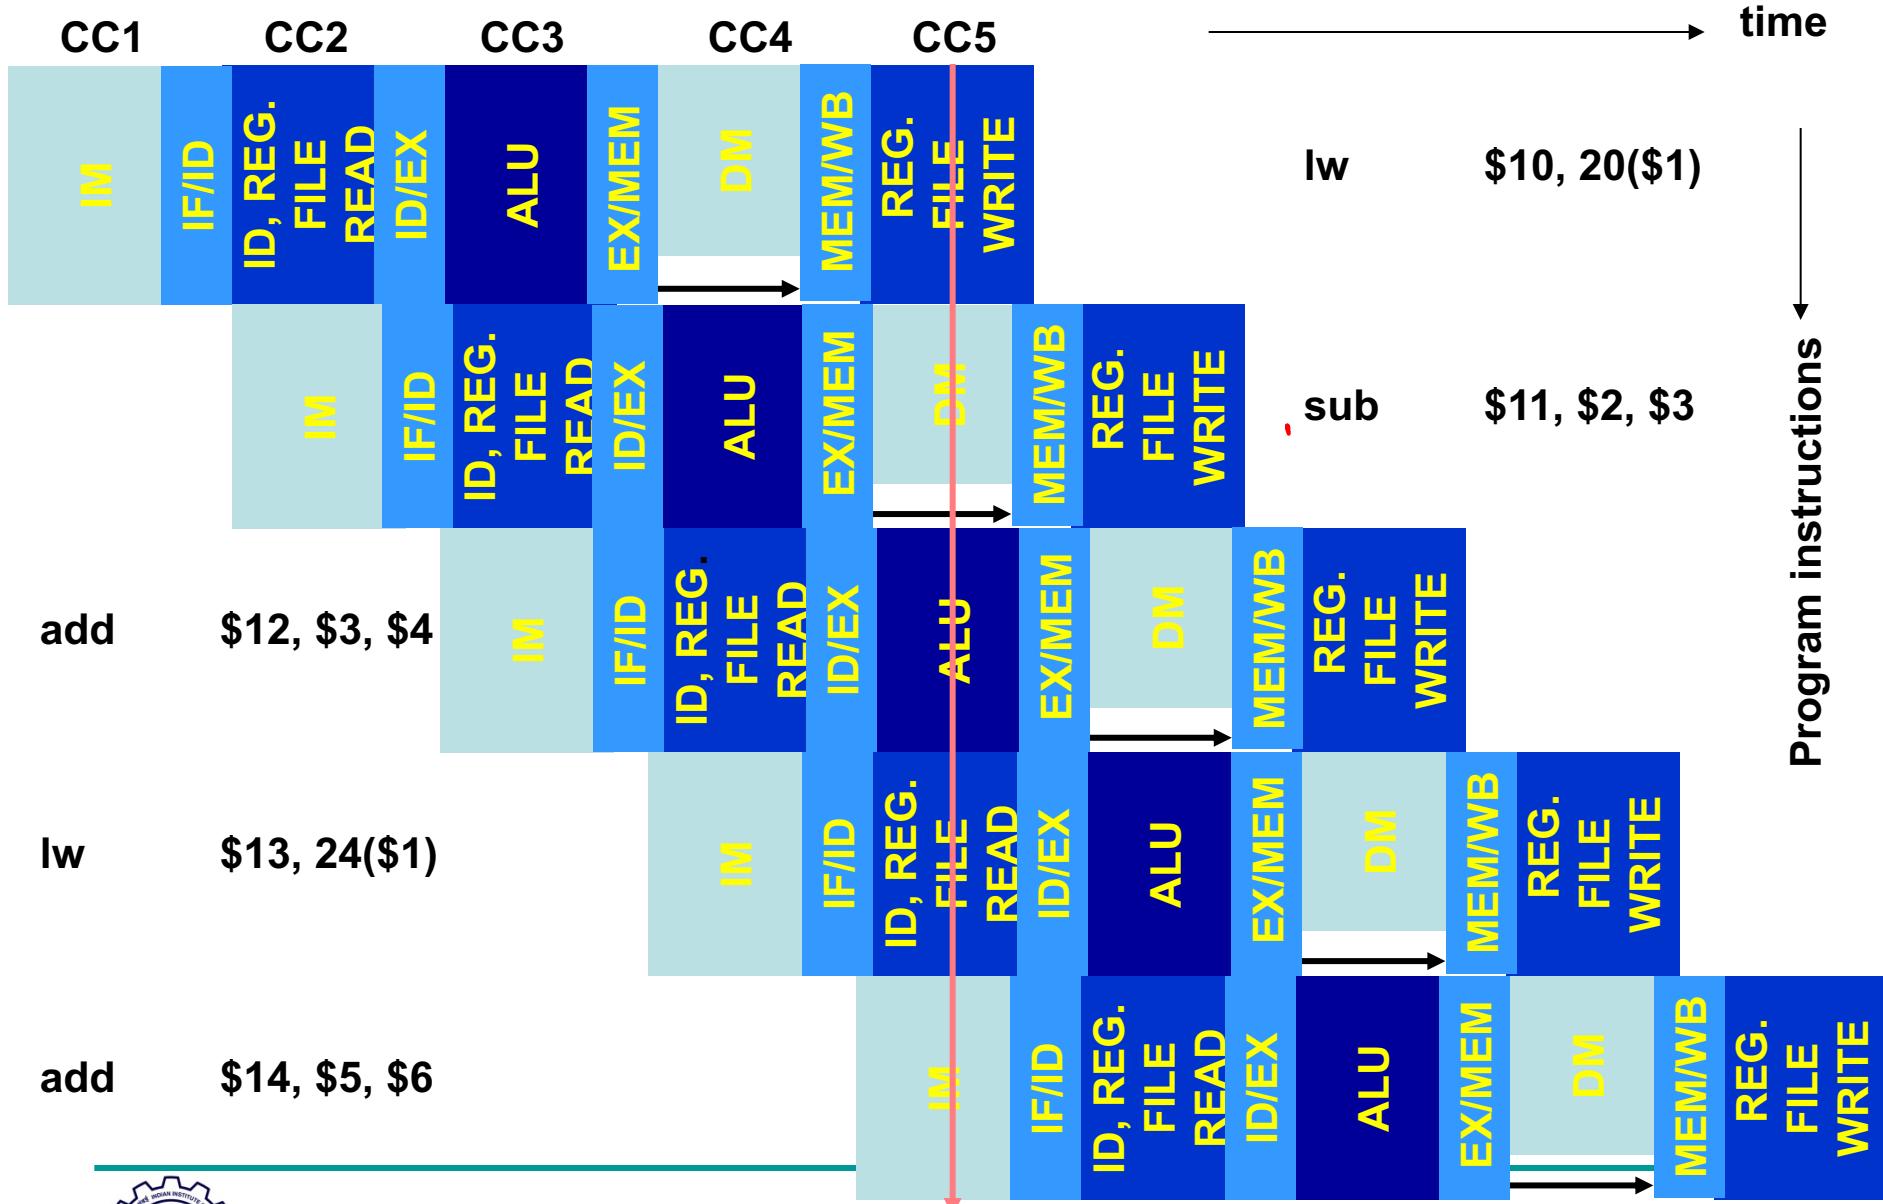
*Although an instruction takes five clock cycles,
one instruction is completed every cycle.*



Pipelined Datapath



Program Execution



	IF	ID	Ex	Mem.	WB/Wdcl -
1.	I_1				G
2	D_2	I_1			
3	I_3	D_2	I_4		
4.	I_4	D_3	D_2	I_1	
5	D_5	D_4	D_3	I_2	I_1
6.	I_6	D_5	I_4	I_3	D_2
7	I_7	I_6	I_5	I_4	D_3 .

latency



Assumption ✓

- 1. All the instructions are independent
- 2. Follow straight code.

↓
No branch. [No change in
control flow)

REACTIVE

- Detect ✓ }
- Correct ✓ }



$$\frac{CPI = 1}{C = \downarrow}$$

T₁
T₂
T₃
f₁
f₂
f₃
f₄
f₅
f₆
f₇



Advantages of Pipeline

- After the fifth cycle (CC5), one instruction is completed each cycle; CPI ≈ 1 , neglecting the initial pipeline latency of 5 cycles.
 - *Pipeline latency is defined as the number of stages in the pipeline, or*
 - *The number of clock cycles after which the first instruction is completed.*
- The clock cycle time is about four times shorter than that of single-cycle datapath and about the same as that of multicycle datapath.
- For multicycle datapath, CPI = 3.
- So, pipelined execution is faster, but . . .



Pipeline Hazards

- Definition: *Hazard in a pipeline is a situation in which the next instruction cannot complete execution one clock cycle after completion of the present instruction.*
- Three types of hazards:
 - Structural hazard (resource conflict)
 - Data hazard ✓
 - Control hazard ✓



Data Hazard

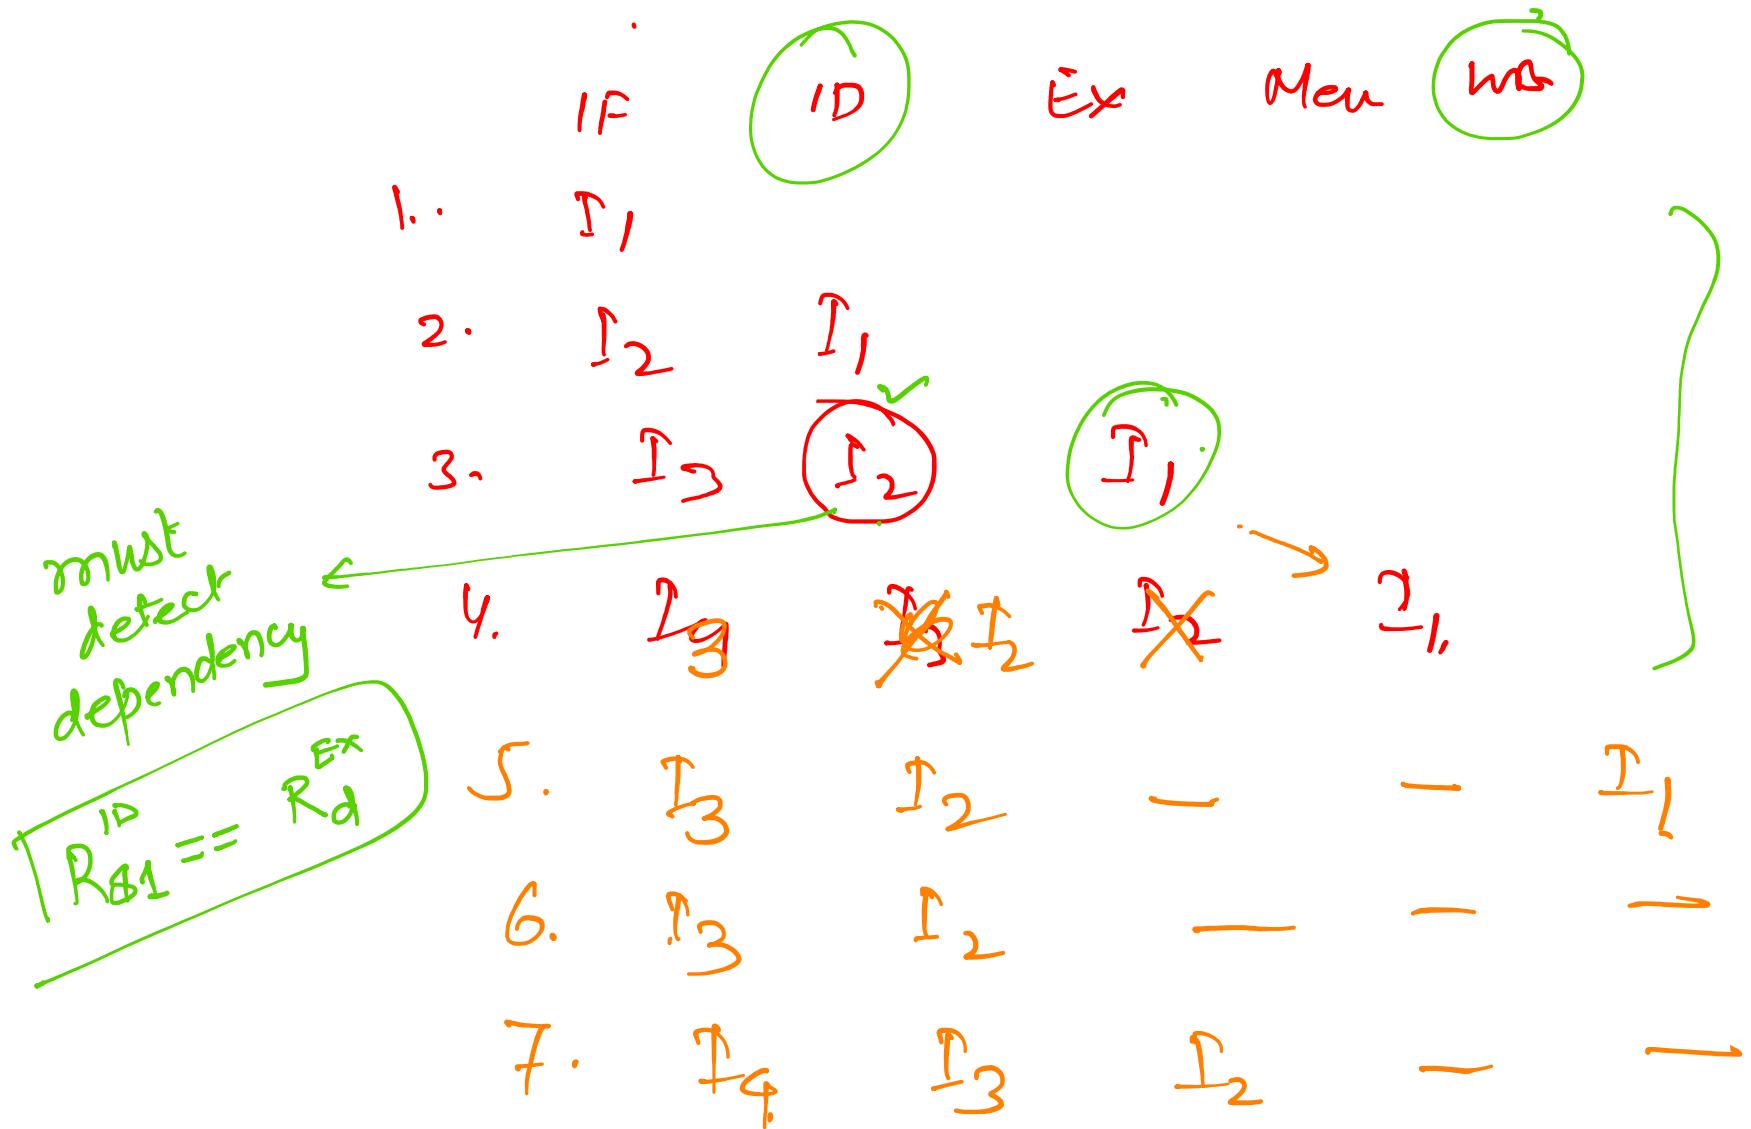
- Data hazard means that an instruction cannot be completed because the needed data, to be generated by another instruction in the pipeline, is not available.
- Example: consider two instructions:

❖ add $\$s0, \$t0, \$t1$ }
❖ sub $\$t2, \$s0, \$t3$ } # needs $\$s0$

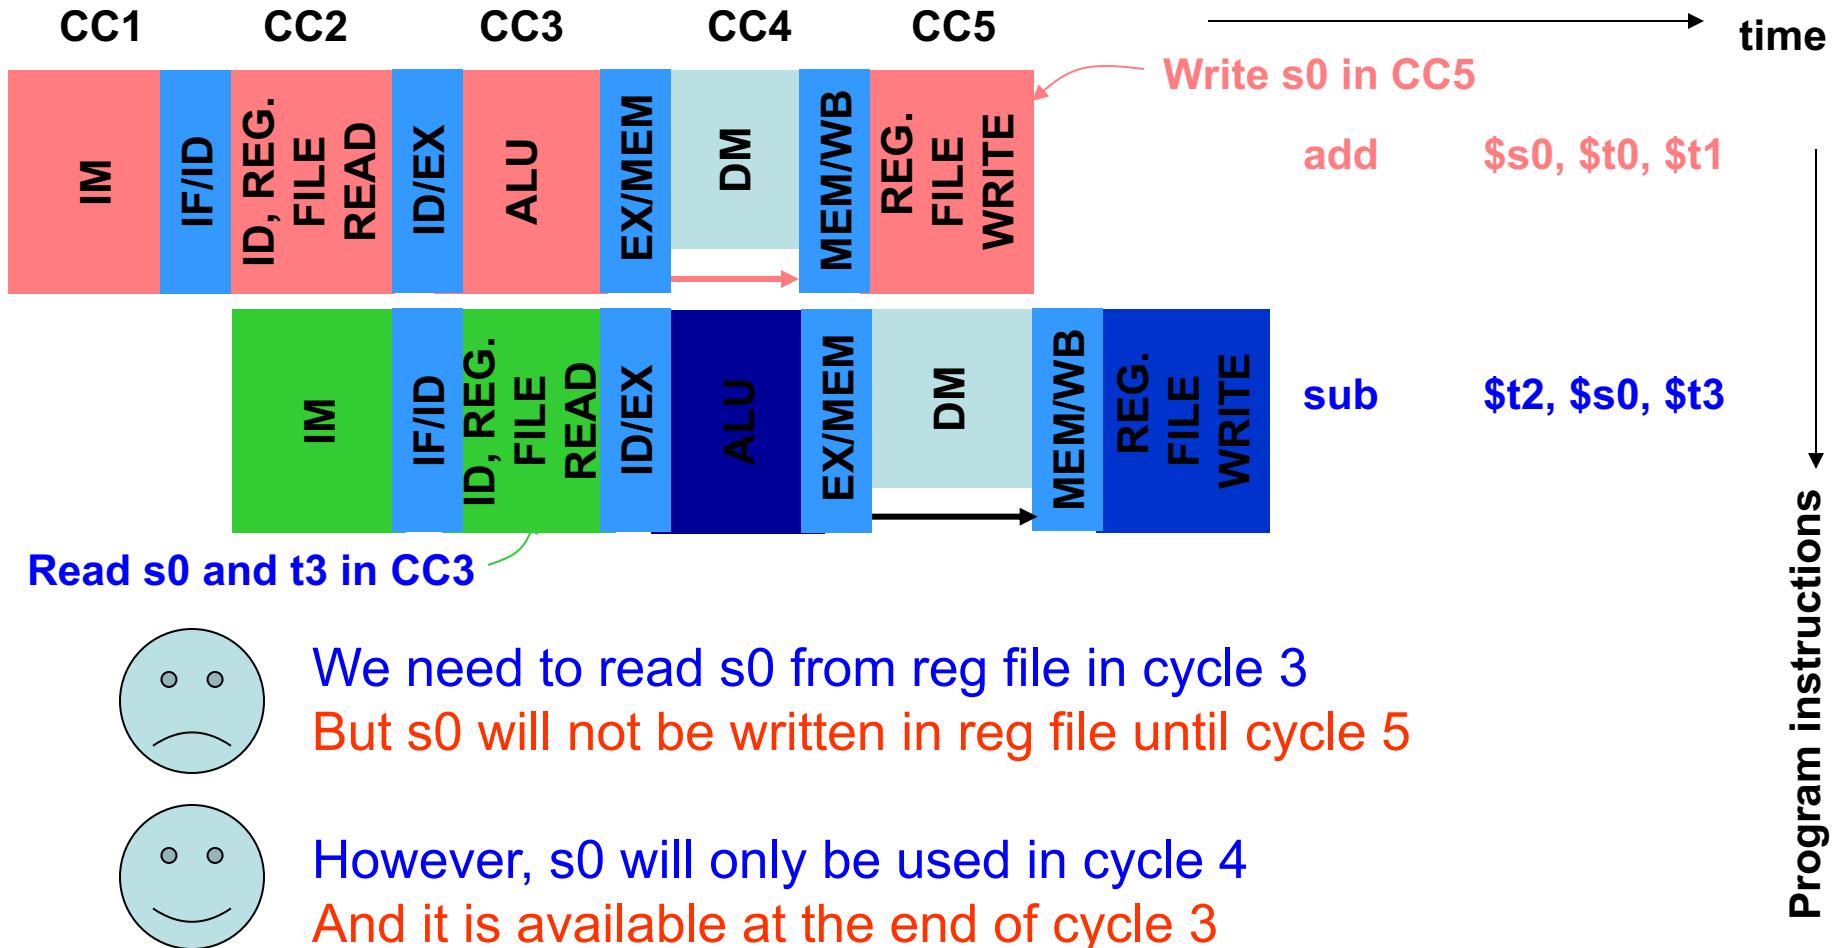
1, add r_1, r_2, r_3)
2, sub r_4, r_1, r_5)

↓
Decide .





Example of Data Hazard

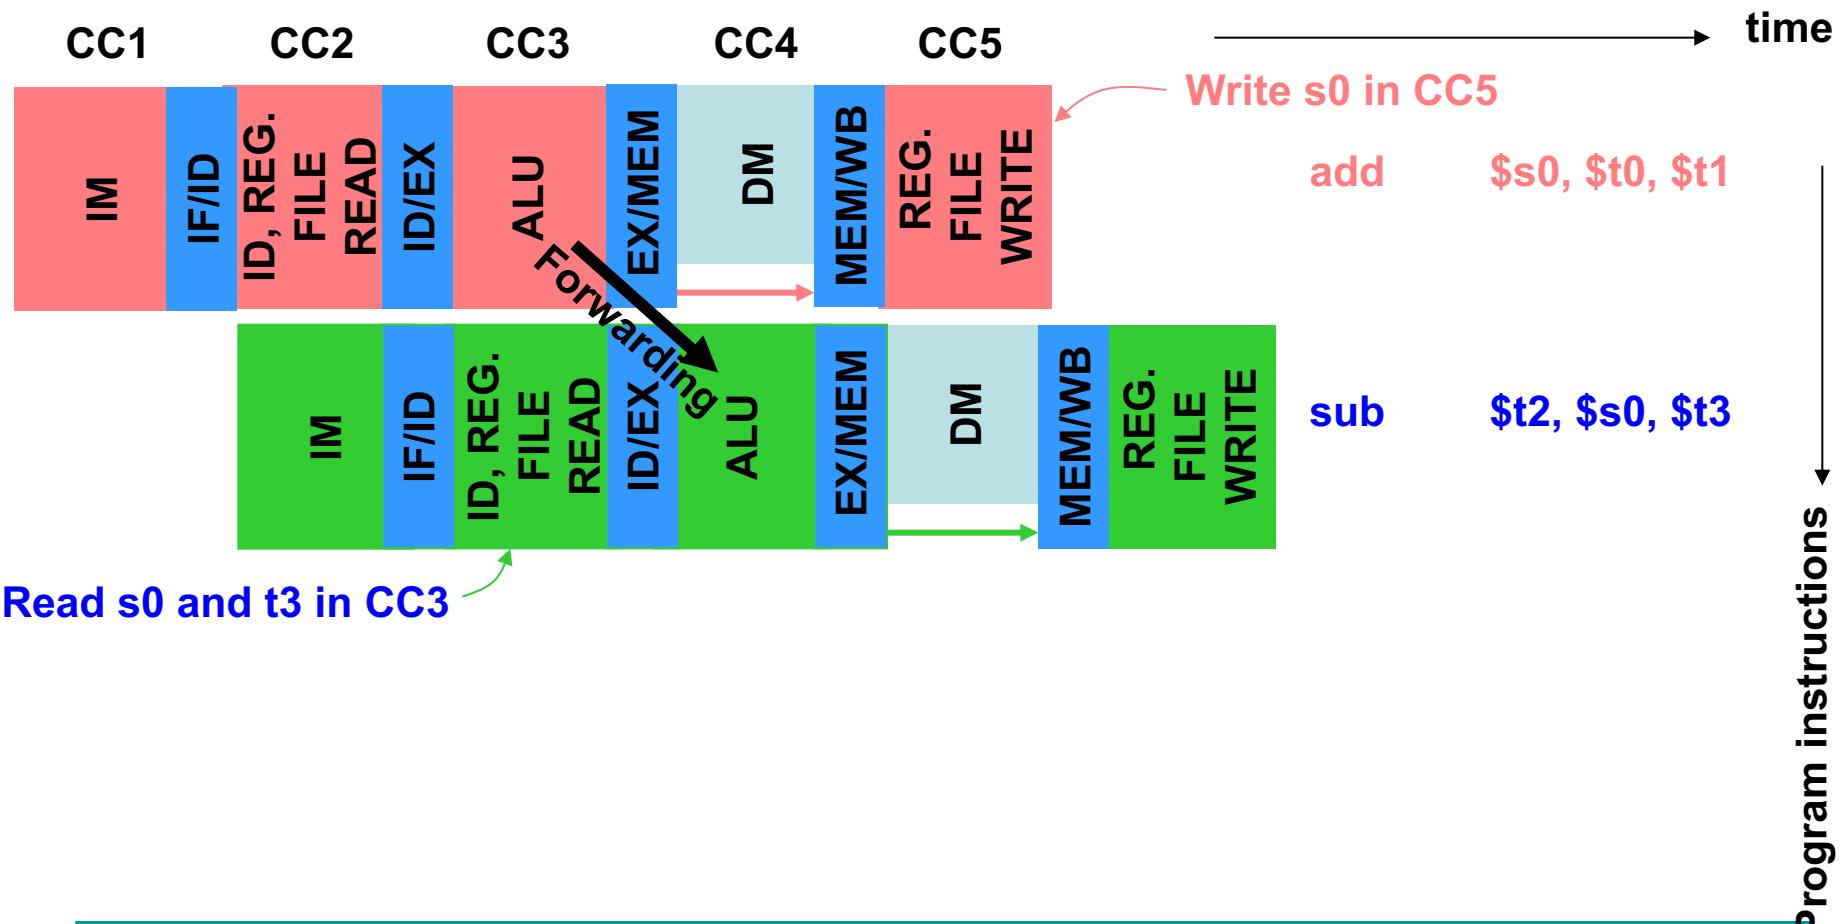


Forwarding or Bypassing

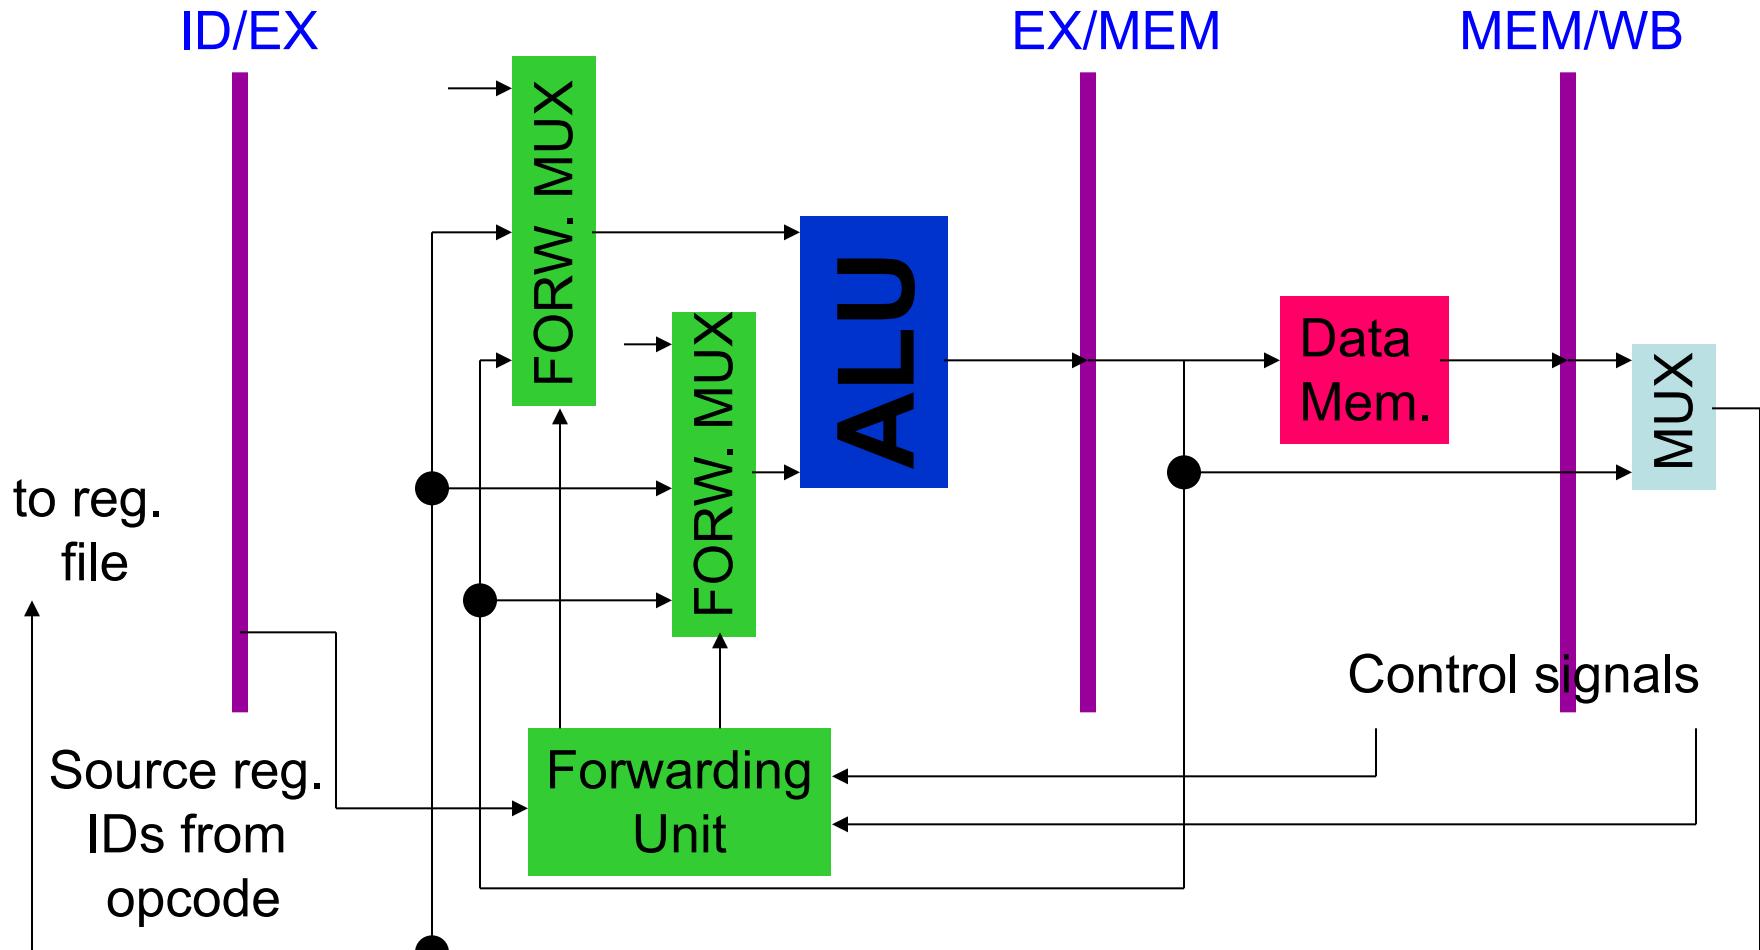
- Output of a resource used by an instruction is forwarded to the input of some resource being used by another instruction.
- Forwarding can eliminate some, but not all, data hazards.



Forwarding for Data Hazard



Forwarding Unit Hardware



Thank You

